

Pegasus 42 Project

The goal for this Gambiconf 2025 workshop is to show how the most basic elements of digital computers and videogame work and how they can be combined to create more complex blocks up to a whole computer.

A simulator (called Digital) will be used to allow the workshop participants to experiment for themselves with all the presented circuits.

1. Switches, Logic Gates, Combinational Circuits

These are the basic blocks of digital circuits.

2. Sequential Circuits

More complex circuits need memory.

3. Processors

mcpu16h is a very simple processor with only 4 instructions and is a good introduction while drv32h is more practical, being compatible with the RISC-V standard.

4. FPGAs and Shin JAMMA

With reconfigurable circuits (FPGAs - Field Programmable Gate Arrays) it is possible to see the simulates circuits working in the real world. As there are many different FPGA boards, the Shin JAMMA standard was created so a single project can work on any of them

5. Video and Audio

Though interesting games can be created to use the text terminal, color graphics output and sound like those from the classic NES (Nintendo Entertainment System or Famicom) allow games that are more fun

6. Pegasus 42

With cache memory and high resolution graphics it becomes possible to run the Squeak Smalltalk programming language

A. History

In addition to the material presented in the workshop, a history of the Pegasus project is also included as an appendix.

1. Switches, Logic Gates, Combinational Circuits

The goal of this workshop is to build a retro computer (roughly equivalent to what people would have bought in the early 1990s for their home) called Pegasus 42 and use it to program simple games. The idea is that the project can be completely understood from the lowest level to the overall system.

For this we will initially use a simulator, and then move on to a FPGA (Field Programmable Gate Array - a chip which can be reconfigured to implement any digital circuit). For each level of abstraction there are several different simulators we can use. From a high to low level we have:

Level	Example Simulators
Architecture	QEMU, MAME
Micro-architecture	SPIM, SimpleScalar
Register Transfer	Verilator, ModelSim
Logic Gate	Digital, TkGate
Switches	IRSIM, MOSSIM
Analog Circuits	Spice, Xyce
Components	TCAD, DEVSIM
Physics	Elmer, Matlab

The advantage of using a simulator instead of the real thing is being able to see details that would be very hard, if not impossible, to measure in the actual circuit. The low level simulators show far more details than the high level ones, but are proportionally slower when running on the same computer. So while it would be possible to simulate a whole computer using Spice, it might take minutes for the simulated circuit to execute a single instruction. We might have to wait for weeks to see if it correctly boots or not. Normally we use the low level simulators for small subcircuits and then higher level simulators for the whole system.

When we said that the goal was to understand Pegasus 42 at the lowest level we exaggerated a bit. We will consider the switch level as being the lowest one in this workshop. Even though *Digital* is not optimized for this level, it is sufficient to illustrate the ideas that will be presented. It is also not optimized for the higher levels, mas for the reduced projects that will be studied it is sufficient (though too slow to show the operation of circuits that output video in a usable manner).

An abstract representation of a system is a box with a number of inputs and some outputs. Normally we will show the inputs coming from the left and the outputs going right, but we can ignore this rule if it makes the drawing more confusing.



Figure 1: system

Digital ou Analog

The first choice we need to make is the nature of the inputs and outputs of our system. For the analog inputs and outputs some quantity in our circuit (voltage, current, etc) is analog to some quantity in the world (temperature, brightness, etc). For the digital inputs and outputs several quantities in the circuit represent a single value in the world. This set of quantities can use separate inputs and outputs (a parallel representation) or a single input ou output over time (a serial representation).

characteristic	Analog	Digital
number of circuits	one	one per digit
precision	depends on circuit quality	always equal to the number of digits
noise	accumulates at each operation	does not pass from input to output

Analog circuits dominated computing until the middle of the 20th century, and telecommunications until the end of the 20th century. The most important factor was the number of circuits since the components were very expensive and their connection was a manual process. With the evolution of integrated circuits the cost became extremely low and the other factors lead to the digitalizations of technology. Our project is digital.

Humanity has used several different digital systems to represent numbers, with the most popular the positional decimal system wth hindu-arabic digits. The

more values each digit can have, the more sensitive to noise it becomes. But the fewer values each digit can have, the more digits are necessary to represent the same number. The best possible protection against noise is when each digit can have only 2 values, like in the positional binary system.

Though the binary system needs more digits (and so, more circuits) than the alternatives, each circuit is simpler so that is the option we shall use.

Logic Gates

Combinational circuits are those whose output (or outputs) depends only on the combination of the inputs. In the case of binary, each digit can only be either 0 or 1. There are several areas of mathematics which are equivalent when only two values are used.

Area					
Boolean Algebra	1	0	inversion	addition	product
Predicate Logic	true	false	not	or	and
Set Theory	universal	empty	complement	union	intersection
	set	set			
Switch Circuits	5V	0V	normally closed	parallel	series

Notations from all of these areas can be used to represent combinational circuits. Yet another possible representation is simply a table with a line for each combination of input values and the corresponding output. We call this a “truth table” even with the values shown are 0 and 1 instead of false and true.

We will not be completely consistent and might talk about a circuit being the in form of “a sum of products” (Boolean Algebra) and another circuit of using “not and” (Predicate Logic). The latter case is why the basic circuits are known as “logic gates”.

To illustrate these ideas we will use Digital, the simulator we had previously mentioned. *Digital* was written in Java, so it is necessary to first install this language on your computer. The advantage of this is that *Digital* runs on computers with different operating systems and different processors. The indicated site is where the source code is, but that is only needed by those wanting to modify the simulator. That page includes a “Download” button which will fetch *Digital.zip* with the most recent version of the tool.

Switches

Since we talked about circuits with switches, lets start there connecting two switches in parallel between a lamp and a power supply. In the “File” menu we select “New”. Using the “Components” menu with “Switches” and “Switch” we

can position two simple switches anywhere we want. Then “Components”, “IO”, “LED” will give us a reasonable approximation of the lamp we wanted (*Digital* has fancier options, but we won’t use them here). Finally in “Components”, “Wires”, “Supply voltage” gets us the power supply we need for our circuit. Note that all circuits need both a power supply and a ground wire, but we normally don’t show those and the simulator works just the same. But if we want to actually build the circuit we need to remember them.

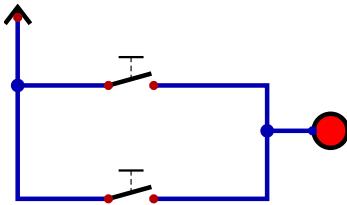


Figure 2: parallel switches

If we imagine two gardens connected using two gates, with one beside the other (in parallel), if one *or* the other is open we can go from one garden to the other. If we simulate this circuit (menu “Simulation”, “Start simulation” on the button with the simple triangle pointing right) we will see that the LED is off. But if we process the switch on top *or* the switch on the bottom (or both) it turns on.

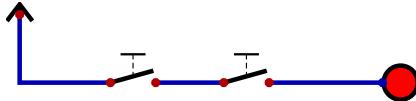


Figure 3: serial switches

If we imagine two gardens connected using two gates, with one after the other with a narrow path between them (in series), it is not enough for one of them to be open. It will only be possible to pass from one garden to the other if the first gate *and* the second gate are open. In this second circuit we connected the switches in series and in the simulation we can see that the LED stays off unless the first *and* the second switch have been pressed.

We have seen two of the three equivalencies between switches and areas of mathematics. The table indicates the remaining equivalency (inversion, not, complement) as being a normally close switch. This kind of switch opens the circuit when pressed. But here we will show an alternative that depends on the device level (the only time we will go down to this level in this project).

We only need one switch here, and in place of the supply we use “Components”, “Wires”, “Pull-Up Resistor”. We also need “Components”, “Wires”, “Ground”. When the switch is open, a current goes through the resistor and into the LED, which shines. As the switch is pressed it offers a path to 0V that draws the current from the resistor instead of the LED, which turns off.

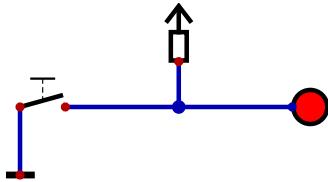


Figure 4: inversion with a switch

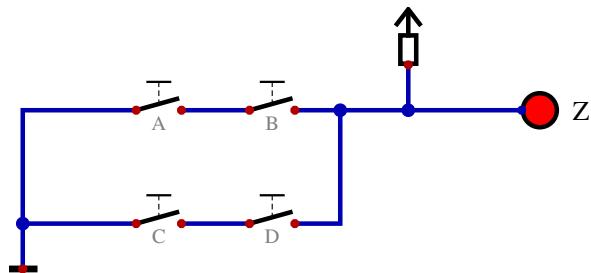


Figure 5: and, or, invert with switches

Here we have a more complex example using the same idea. Normally the AOI gate(and/or/invert) is not considered to be a basic logic gate and we will not see it again in this project, but it is sufficiently useful that it is often included in a library for integrated circuit design. In Boolean algebra we have:

$$Z = !(Ax_B + CxD)$$

while in predicate logic it would be:

$$Z = \text{not}((A \text{ and } B) \text{ or } (C \text{ and } D))$$

A fourth representation of the AOI circuit (with the first being the figure, or schematic, the second the boolean algebra and the third the logic equation) would be the truth table:

A	B	C	D	Z
open	open	open	open	on
open	open	open	closed	on
open	open	closed	open	on
open	open	closed	closed	off
open	closed	open	open	on
open	closed	open	closed	on
open	closed	closed	open	on
open	closed	closed	closed	off
closed	open	open	open	on
closed	open	open	closed	on
closed	open	closed	open	on

A	B	C	D	Z
closed	open	closed	closed	off
closed	closed	open	open	off
closed	closed	open	closed	off
closed	closed	closed	open	off
closed	closed	closed	closed	off

The table shows a problem - the input have one nature (they are controlled by a human finger) while the output is of a different nature (light coming from the LED). If we want the outputs of a circuit to be used as inputs of another circuits in order to build larger systems, they need to be of the same kind. Fortunately switches which are controlled using electricity have been invented: relays (1835), vacuum tubes (1904) and transistors (1947). Even though *Digital* can simulate relays in a limited way (like its switches) we will replace the switches in the AOI circuit with MOSFET (Metal/Oxide/Silicon Field Effect Transistors) of the N type (negative) which is used to make integrated circuits.

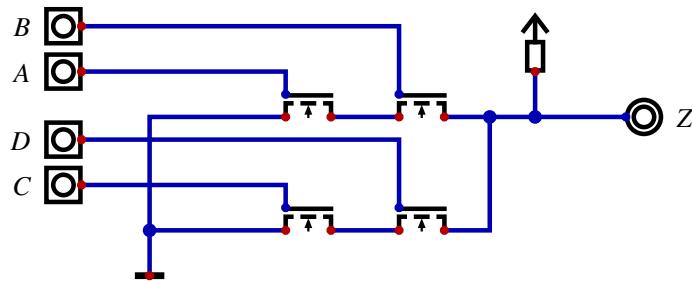


Figure 6: and, or, invert NMOS

Besides replacing the switches with “Components”,“Switches”,“N-Channel FET” we also use “Components”,“IO”,“Input” as well as “Output” from the same menu to show that these signals can come from another circuit and that the result can go to another circuit. While simulating we can change the values of the inputs and observe the value of the output.

In the “Analysis” menu, the “Analysis” item will create a truth table for the circuit and we can verify that it is the same one as for the circuit with switches. A problem with this type of circuit, which we call NMOS, is that whenever the output is 0 there is a current going through the resistor and generating heat (and draining the battery if that is where the power is coming from). Another kind of transistor, the “P-Channel FET”, is the opposite of the N type and conducts current when the input is 0. If we replace the resistor with a complementary circuit of the N transistors using the P transistors (placing them in series where the other is parallel) the circuit will work just the same but without a constant current.

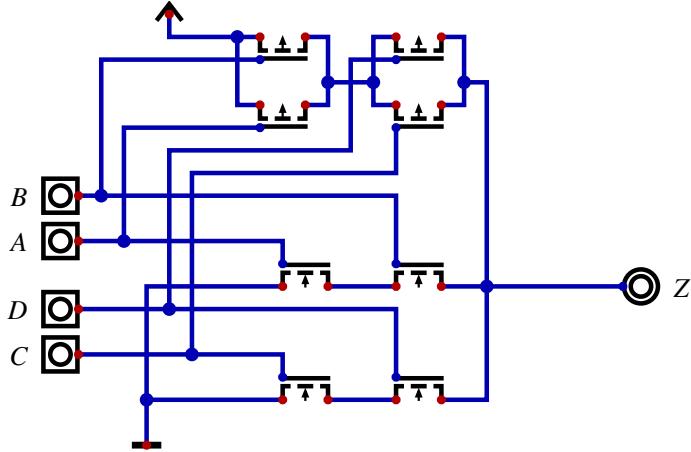


Figure 7: and, or, invert CMOS

The extra complexity of CMOS relegated it to niche applications (like digital watches) in the 1960s and 1970s, but it replaced nearly all other kinds of circuits in the 1980s when the increasing number of transistors per chip (Moore's Law) made having twice as many transistors worth it to reduce power.

one input Logic Gates

So far we have seen a single logic gate with one input: Not.

With a single input, only two combination of inputs are possible: either 0 or 1. Its truth table will be two lines in size. The output in each line can have two values, so there are $2^2 = 4$ possible truth tables.

A	Z
0	0
1	0

In the first gate the output is always 0. It isn't surprising that we didn't talk about it. In circuit terms we only need to connect the output to the ground wire.

A	Z
0	1
1	0

This is the Not gate that we have already seen.

$$\begin{array}{r} \overline{\overline{A}} \quad Z \\ \hline 0 \quad 0 \\ 1 \quad 1 \\ \hline \end{array}$$

Here the output is the same as the input. Just like in the first circuit we can implement this with just a wire.

$$\begin{array}{r} \overline{\overline{A}} \quad Z \\ \hline 0 \quad 1 \\ 1 \quad 1 \\ \hline \end{array}$$

The final gate is an output that is always one. This can also be done with a wire connected to the power supply.

So out of the 4 possible gates, only Not is interesting. Note that if we consider the value of Z in each table as the bits of a binary number with the first line being the least significant digit, we can call these gates “gate 0” ($Z = 0$), “gate 1” ($Z = !A$), “gate 2” ($Z = A$) and “gate 3” ($Z = 1$).

two input Logic Gates

Using the same thinking, a two input gate has 4 possible input combinations and therefore a truth table with 4 lines. Using the same scheme to number them we will have a 4 bit binary number indicating that there are $2^4 = 15$ possible logic gates.

outputs	equation	name
0 0 0 0	$Z = 0$	
0 0 0 1	$Z = !(A+B)$	NOR
0 0 1 0	$Z = Ax!B$	
0 0 1 1	$Z = !B$	
0 1 0 0	$Z = !AxB$	
0 1 0 1	$Z = !A$	
0 1 1 0	$Z = !(AxB)+(Ax!B)$	XOR
0 1 1 1	$Z = !(AxB)$	NAND
1 0 0 0	$Z = AxB$	AND
1 0 0 1	$Z = (AxB)+(!Ax!B)$	XNOR
1 0 1 0	$Z = A$	
1 0 1 1	$Z = A+!B$	
1 1 0 0	$Z = B$	
1 1 0 1	$Z = !A+B$	
1 1 1 0	$Z = A+B$	OR

outputs	equation	name
1 1 1 1	$Z = 1$	

Gates 0000 and 1111 don't actually have any inputs, while 0011, 0101, 1010 and 1100 ignore one of the inputs. They are actually the gates we already saw above.

6 gates have names in the menu "Componentes", "Logic" as well as a corresponding drawing. There is a special shape for AND (1000), OR (1110) and XOR (exclusive OR - 0110) and for their inverses we just add a little circle to the output and an "N" to the beginning of the name. *Digital* also allows a little circle to be added to any of the inputs which means we can use an AND for gates 0010 and 0100 and an OR for 1011 and 1101.

Sum of Products If we look at the line for XOR and XNOR we will notice that they are the most complicated ones. In the case of XOR the first product (AND) is !AxB which corresponds directly to the left 1, while Ax!B is what generated the right 1.

outputs	product
0 0 0 1	!Ax!B
0 0 1 0	Ax!B
0 1 0 0	!AxB
1 0 0 0	AxB

So the XOR is the sum (OR) of the second and third products of this table. This is actually true for any logic gate and can be expanded for any number of inputs. This means that nobody will ever discover a new logic gate in the future that we don't know how to implement. The truth table can be transformed directly into a circuit.

That doesn't mean that the circuit created this way will be very good. Using this method for the next to the last logic gate we would have

$$Z = (\text{Ax!B}) + (\text{!AxB}) + (\text{AxB})$$

But we know that a simple OR gate will do the same thing. Fortunately Boolean algebra has rules for simplification much like the rules in normal algebra, and there is a graphical method (called Karnaugh Map, which is one of the things *Digital* can generate) to reduce the logic to the minimum while keeping the same operation.

Multiple Bits

One disadvantage of digital circuits relative to analog ones is the repetition of the same circuit for each digit. If we use 32 bits to represent values, we will have 32 copies of each circuit.

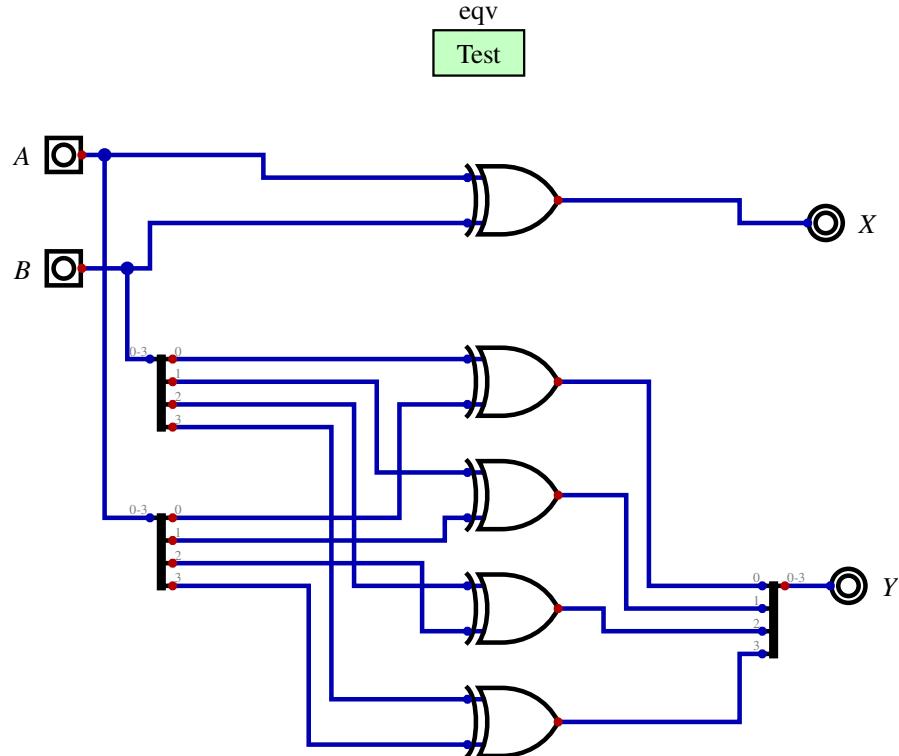


Figure 8: multiple bits

Digital has a feature which can reduce this complexity. For each input and output, besides the name we can define a “width” in number of bits. In the above circuit we changed A , B , X and Y to have 4 bits each. In “Components”, “Wires”, “Splitter/Merger” we have a way to connect signals with different numbers of bits. In the two on the left the input was configured as “4” and the output as “1,1,1,1” while the one on the right was configured as the opposite. In addition, each normal component can be configured to have a given width. The exclusive or gate at the top was configured to 4 bits while the four below it as 1 bit each.

Both circuits should be completely equivalent, but the one on top is easier to understand since it is much smaller. And this is for only 4 bits - the gain for, for example, 32 bits is proportionally greater. While editing the circuit, *Digital* does not give any visual indications that the top gate is different than the others, that

the inputs and outputs are for multiple bits nor even that the wires connecting them will be carrying multiple bits. During simulation, however, the wires with 1 bit are shown in dark green (for 0) or light green (for 1) while those with several bits continue dark blue with the current value shown right above the wire. And clicking on a 1 bit input will invert its value while on a multiple bit input this will open a dialog box to define the new value.

Using “Analysis”, “Analysis” we can see the truth table and compare the bits from X and Y to check that the circuits are actually equivalent. With 256 lines, however, this confirmation is quite tiring. And if we make any changes to the circuit we will have to repeat this careful examination of the truth table. Fortunately, *Digital* allows us to automate this using “Components”, “Misc”, “Test Case”. We edit the test (which we name “eqv”) to:

```
A B X Y
```

```
loop(a,16)
  loop(b,16)
    (a)  (b)  (a^b)  (a~b)
  end loop
end loop
```

Using “Simulation”, “Run tests”, all 256 combinations of a and b are generated for A and B and X and Y are compared with $a \text{ XOR } b$. All tests that pass are shown in green and all tests that fail are shown with a red “x”. It is possible to examine the details to know where a failure happened.

Another way that *Digital* hides complexity is the use of hierarchical projects. Several components can be combined into a circuit that is then shown as a single block in a higher level circuit. Any practical project should make extensive use of this as well as include many tests for each sub-circuit. But since the goal of this workshop is to show the complexity of a computer, the projects that will be shown will be as “flat” as possible. Sub-circuits are often compared with subroutine calls in programming languages but they are actually more like macros.

Decisions

In programming languages we have constructions like “if A then B else C” to use one input to select between two other inputs.

Looking at the truth table we can see that $Z = B$ whenever $A = 1$, but $Z = C$ when A is 0. This means that combinational circuits can make decisions. It is possible to select between more than two alternatives, like in the “switch/case” statements that programming languages have.

This is more complicated to test because with $8+3 = 11$ inputs there are 2048 possible combinations. The complete test ($t8x256$) is possible:

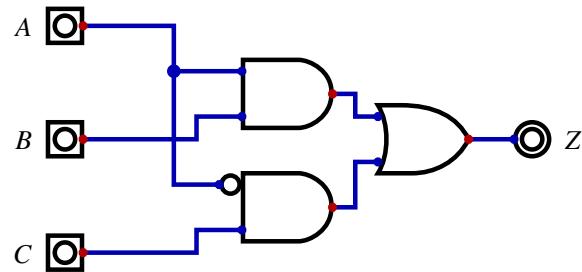


Figure 9: 2 input multiplexer

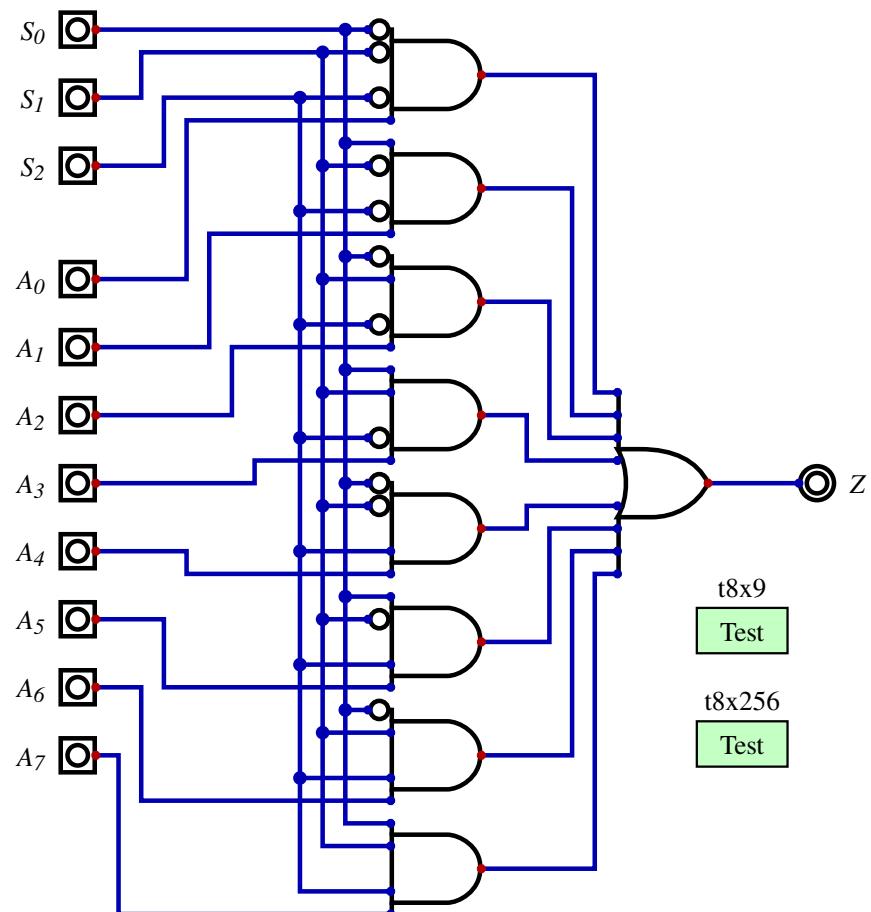


Figure 10: 8 input multiplexer

```

S_2 S_1 S_0    A_7 A_6 A_5 A_4 A_3 A_2 A_1 A_0  Z

loop(s,8)
  loop(a,256)
    bits(3,s) bits(8,a) bits(1,a>>s)
  end loop
end loop

```

Here we are confirming that the output corresponds to the bit selected by S (Z is calculated by right shifting so that the least significant bit is the one from the desired input). A more careful test would be to check that when all inputs are 0 the output is also 0 and then to turn on each input, one at a time. The output should remain at 0 except when it is the selected input that goes to 1. Instead of 8×256 tests we need only 8×9 . In this case the more complete test is actually better, but during the fabrication of some product where part of the cost is the time spent using test equipment the reduced solution would be more interesting.

Our project will use many multiplexers and the circuit above is rather big (and would have to be repeated 32 times to select between 8 values of 32 bits each). At the logic gate level this is the best solution, but if we go down to the switch level we can save transistors.

This circuit passes the same tests as the previous one. In practice this use of NMOS pass transistors reduces the signal level, but adding two inverters to the output will eliminate this problem. Using pairs of NMOS and PMOS pass transistors is another solution, but the wiring becomes a bit more complex. The idea of showing this was so we can use multiplexers in our projects without worrying too much about their cost.

Numbers

A very popular idea is that computers only manipulate numbers but we can interpret these numbers as being letters, colors, sounds, etc. This is subtlety wrong - computers can manipulate representations of many things, including numbers. This is not easy to notice for positive integer numbers, but for negative numbers or floating point numbers this becomes clearer.

The first detail we need to notice is the difference between “+” in Boolean algebra and in binary arithmetic:

A	B	Boolean A+B	Binary A+B
0	0	0	0
0	1	1	1
1	0	1	1
1	1	1	10

The result in the last line has more digits than the inputs in the case of binary



Figure 11: 8 input multiplexer

arithmetic. That is also the case for decimal arithmetic: adding two decimal numbers with 5 digits each might generate a result with 6 digits. And for each digit the result will be from 0 to 18, with the last being 2 digits.

Note that the last line in the table is not saying that “one plus one equals ten”. The same way that the decimal number 307 means $3 \times 100 + 0 \times 10 + 7 \times 1$, the binary number 1010 means $1 \times 8 + 0 \times 4 + 1 \times 2 + 0 \times 1$ which is the number ten. In the decimal positional system the rightmost digit is the unity and each digit to the left is worth ten times as much. In the binary positional system the rightmost digit is the unity and each digit to the left is worth two times as much. The binary number 10 is $1 \times 2 + 0 \times 1$ which is two.

We shall call the digits from the addition of binary 1 bit numbers S (from “sum”) for the least significant and C (for “carry”) for the most significant.

A	B	C	S
0	0	0	0
0	1	0	1
1	0	0	1
1	1	1	0

C has the truth table of the AND logic gate and S of the Exclusive OR gate. So an adder is:

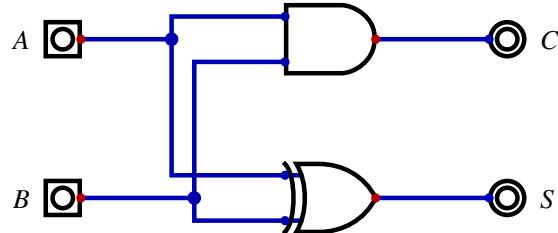


Figure 12: half adder

For the least significant digits of the input numbers this does the job. But for any other digits we need to take into account the “carry” from the digit immediately to the right. That is why we call this circuit a “half adder” and use two of them to create a “full adder”, which is a circuit with 3 inputs and 2 outputs.

Repeating the full adder 4 times we can add two binary numbers with 4 bits each. Actually, we could have used a half adder for the least significant digit but doing it this way it becomes easy to combine several of these blocks to handle larger numbers.

We can test that we are actually adding numbers:

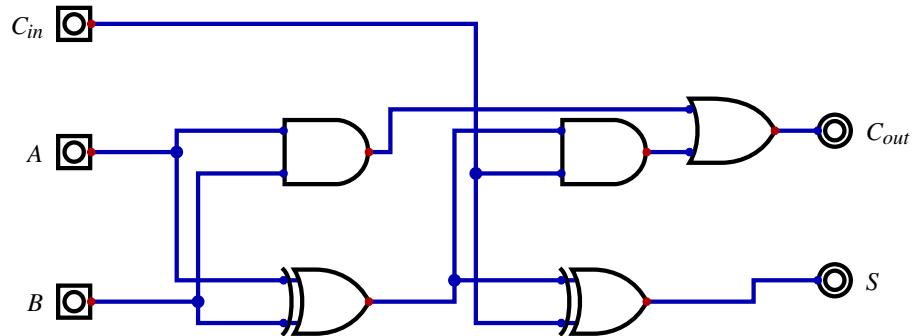


Figure 13: full adder

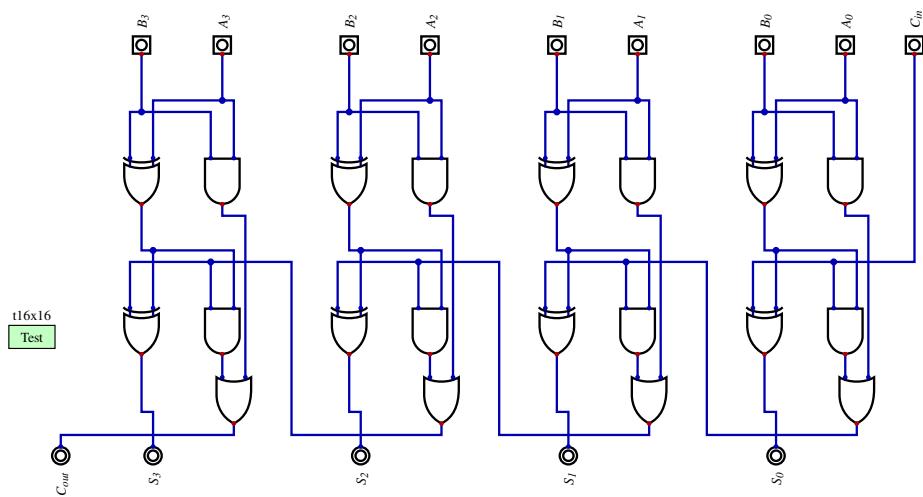


Figure 14: 4 bit adder

```

C_in  A_3 A_2 A_1 A_0  B_3 B_2 B_1 B_0  C_out S_3 S_2 S_1 S_0

loop(a,16)
  loop(b,16)
    0 bits(4,a) bits(4,b) bits(5,a+b)
    1 bits(4,a) bits(4,b) bits(5,a+b+1)
  end loop
end loop

```

Note that with C_{out} and S_3 to S_0 the result is 5 bits long, one more digit than the inputs.

This “ripple carry adder” is not very fast. The carry goes from one digit to the next and only arrives at the most significant digit with a long delay. If we double the number of digits the adder will operate at half the speed. There are more complex circuits that reduce this problem, but for our project this simple solution is good enough.

For subtraction we face a new problem: negative numbers. We mentioned previously that computers manipulate representations and not numbers, but it is easy to ignore the difference in the case of positive numbers. For negative numbers we must choose between several representations, each with its advantages and complications.

The most popular representation for decimal numbers is the sign/magnitude one. A special digit to the left indicate if the number is positive (“+” or nothing) or negative (“-”). The remaining digits indicate the number’s absolute value. The same idea can be applied to binary numbers, and we can even use 0 and 1 to indicate positive and negative numbers to make the system even more uniform. Two problems with this system is the existence of two different zeros (+0 and -1) and the fact that addition and subtraction are slightly different operations and it is necessary to examine the signs of the operands and compare their magnitudes to select the right operation.

A representation that was only every used in a few of the very first mechanical calculators was nine’s complement. A negative number is represented by replacing each digit by nine minus that digit. The negative of 307, for example, would be 692 (and in the case of a calculator with six digits 000307 negated would be 999692 which makes it easy to interpret the leftmost digit as indicating the sign). Now subtractions is just inverting the second operand and adding (ignoring any carry generated during that operation): $000531 - 000307 = 000531 + 999692 = (1)000223$. Except this is not the right answer since $307+223 = 530$. This is one of the problems with nine’s complement: we need to add one to correct the result. And it also has two zeros. In the binary case the equivalent is one’s complement (which is the NOT function).

Ten’s complement is a slight modification of the previous scheme that corrects both problems. To negate a number we subtract each digit from nine and then we add one. The ten’s complement of 000307 is 999693 and now additions will

directly give the right answer. Besides that 000000 continues to represent zero but 999999 now means negative one - there no longer is a negative zero. One detail is that there are more negative than positive numbers for a given digit length, but that is only an aesthetic issue. The equivalent binary system is the two's complement.

The last system we will look at is the bias. Here we add a value to all numbers so they all are positive. Selecting the bias as half of the largest possible number plus one, the addition of two number will convert the two biases into a carry, and then we need to add a third bias to adjust the answer. 531 is represented by 500531 and -307 by 499693. $500531 + 499693 + 500000 = (1)500224$.

Comparing these representations in the case of 3 digit binary numbers (with the values shown in sign/magnitude decimal notation):

representação	000	001	010	011	100	101	110	111
positive	0	1	2	3	4	5	6	7
sign/magnitude	+0	+1	+2	+3	-0	-1	-2	-3
one's complement	+0	+1	+2	+3	-3	-2	-1	-0
two's complement	+0	+1	+2	+3	-4	-3	-2	-1
bias	-4	-3	-2	-1	+0	+1	+2	+3

All these representations were used in the history of computing, but in the 1960s two's complement became dominant and it is what we will use in our project. In the IEEE 754 standard for floating point numbers, however, the mantissa uses the sign/magnitude representation while the exponent uses the bias representation.

With a small change to the 4 bit adder it can also subtract using the two's complement representation. We said that NOT can generate the one's complement, but using a XOR for each digit we can control whether we do this inversion or not. In the -16×16 subtraction test we see that we are actually calculating $a-b-1$ which is a problem we saw with one's complement. But using C_{in} we can correct that, getting the two's complement. In the subtract mode the circuit generates an inverted C_{out} and the test takes this into account.

```

invB  C_in   A_3 A_2 A_1 A_0   B_3 B_2 B_1 B_0   C_out S_3 S_2 S_1 S_0

loop(a,16)
  loop(b,16)
    1 0 bits(4,a) bits(4,b) bits(5,(a-b-1)^16)
    1 1 bits(4,a) bits(4,b) bits(5,(a-b)^16)
  end loop
end loop

```

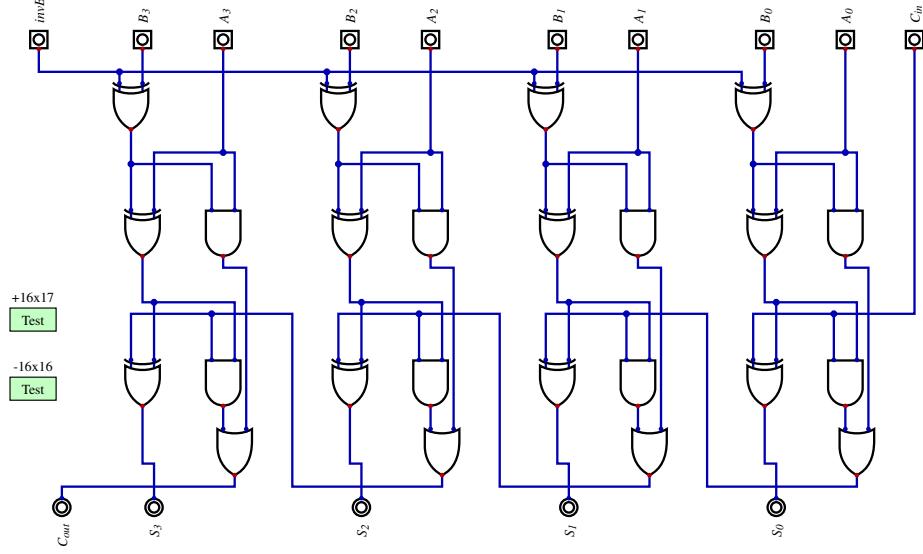


Figure 15: 4 bit adder and subtractor

2. Sequential Circuits

Time is not a factor for a combinational circuito. Of course, given that it is not infinitely fast the answer will only be ready after a certain delay after the inputs receive their values.

In contrast, time is fundamental for sequential circuits. The inputs arrive as a sequence over time using the same signals and the results are also sent as a sequence of values over time. We can convert the abstract combinational circuit we previous saw into a sequential system by connecting some of its outputs to inputs. We call the value being fed back the “current state” of the system.

In practice a circuit like this can even work, but it would be rather unstable since some paths generating some of the bits of the current state might have longer delays than those for other bits. The most popular solution is to add a “clock” signal so the feedback can happen in a more controlled fashion.

Oscillators

Digital refuses to simulate this circuito, which is the simplest sequential system. One problem is that the circuit is contradictory: the output would have to be 1 for a 0 input, or 0 for a 1 input. But there is a wire connecting them so they should have the same value. If we actually build this circuit it will keep quickly alternating between 0 and 1. This is what we call an “oscillator”. The oscillation frequency depends on the speed of the inverter and the delay of the wire. The main problem from a simulation viewpoint is that the initial value is unknown,

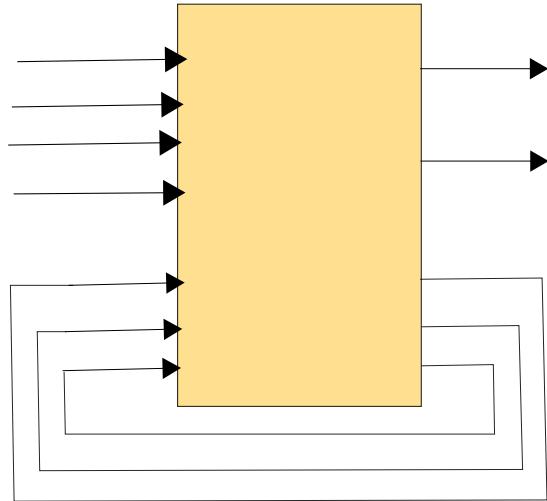


Figure 16: sequential system

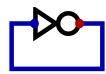


Figure 17: one inverter oscillator

which we can solve with an initialization signal (we will call it “reset”).

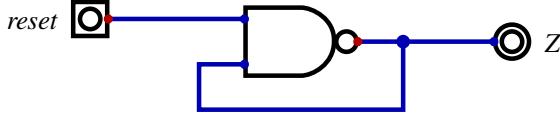


Figure 18: one nand oscillator

If we try to simulate this circuit it seems to work, but as soon as we change *reset* to 1 an error occurs. The solution is to simulate step by step. After changing *reset* to 1, at each step in the simulation the output alternates between 0 and 1.

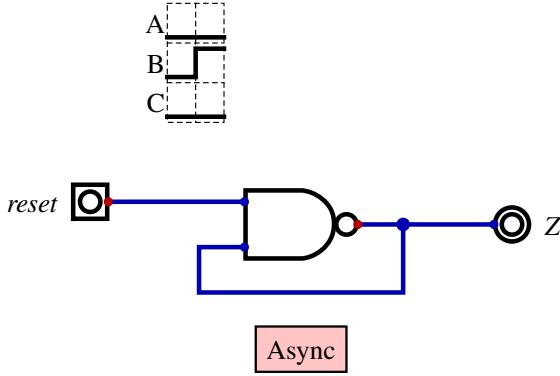


Figure 19: one nand oscillator

Adding “Components”, “Misc”, “Async” changes how the simulation’s visualization happens, and with “Componentes”, “IO”, “Graph” we can visualize the input and outputs over time, which is to sequential circuits what truth tables is to combinational circuits.

Memory

Connecting two such oscillators in series makes it stop oscillating since an even number of inversions is not contradictory. This circuito is “bi-stable”, meaning there are two different situations in which the circuit stays put. The output of the first part is always the opposite of the output of the second part, so we use a “!” in front of its name. We also do that for the inputs to indicate that these should normally stay at 1 and should only go to 0 when we want it to do its job.

The name of this circuito is “flip-flop” to indicate its bi-stable nature, and it is a computer’s simplest memmory capable of holding 1 bit. A more complete name is “RS flip-flop” since the inputs *!reset* and *!set** are separate. The circuit “remembers” that a negative pulse has arrived at *!set* even after that pulse is no longer there. Additional pulses don’t do anything. The same thing for pulses at *!reset*.

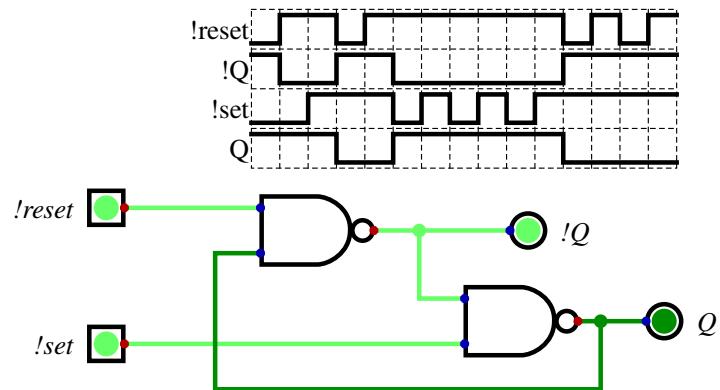


Figure 20: flip-flop with nands

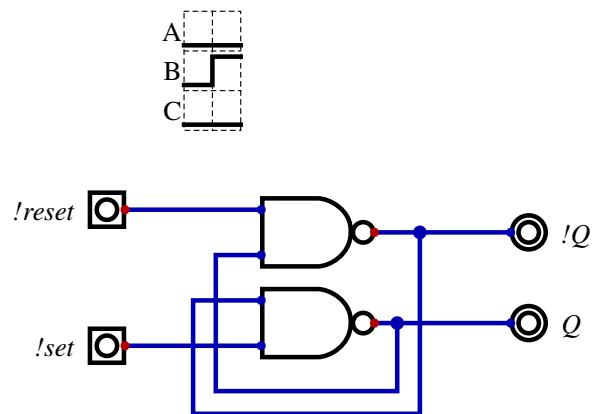


Figure 21: flip-flop with nands

This is exactly the same circuit, but drawn in the style of the abstract sequential system with the outputs going all the way around the circuit to connect to the inputs. The idea is that even with the complications of the next circuits we can still keep in mind that this feedback is happening.

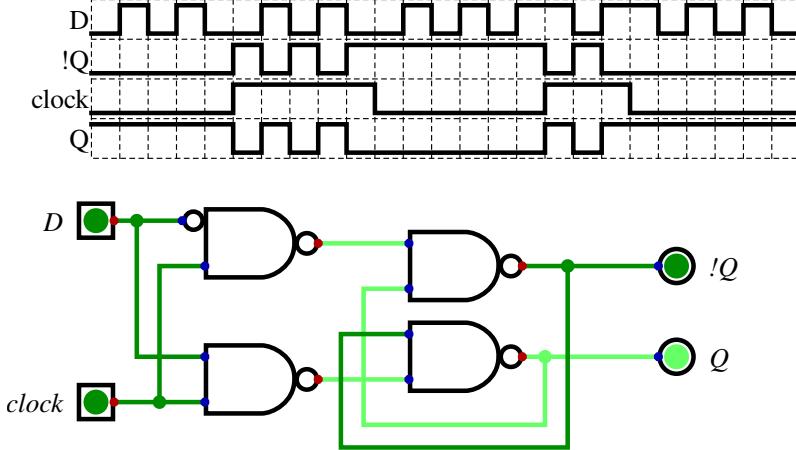


Figure 22: D flip-flop

Instead of directly generating *!reset* and *!set* it is more convenient to have a data signal *D* and a *clock* signal. When the clock is high, the output reflects *D* with a small delay. When the clock is low the output value doesn't change no matter what is happening in *D*.

We can also use a multiplexer to have the same functionality. Though in this case we don't have the inverted input. Modern project avoid using latches because during half of the clock cycle we can't trust the output. Using two latches in a row operating on opposite levels of the clock we can have an output that remains stable during nearly the whole cycle, with the only uncertainty very close to one of the clock edges.

We can have this same functionality by selecting “Components”, “Memory”, “Register”. But it is important to understand what is happening inside the component. We also replaced the *clock* input from a regular one to a special input. In this simulation it didn't make any difference, but we can configure it to pulse at a give rate without having to keep manually clicking on it. And in the tests this kind of input can be set to not only 0 and 1 but also C to indicate a rising edge.

Finite State Machines

Using registers we can solve the stability problem mentioned for the abstract sequential circuit. We just have to pass the output signals to be fed back through registers and make the register outputs be the ones to actually go back into the inputs. This way it doesn't matter if different bits have different delays since all

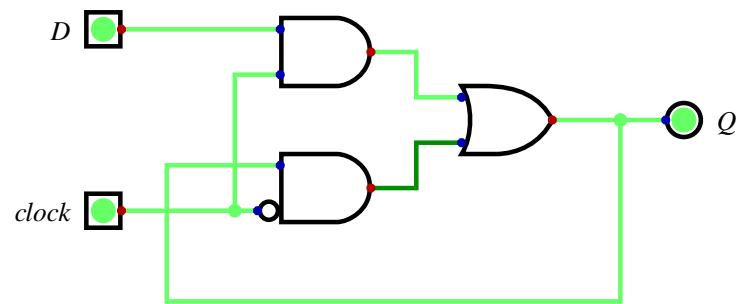
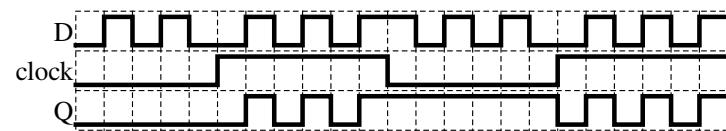


Figure 23: latch

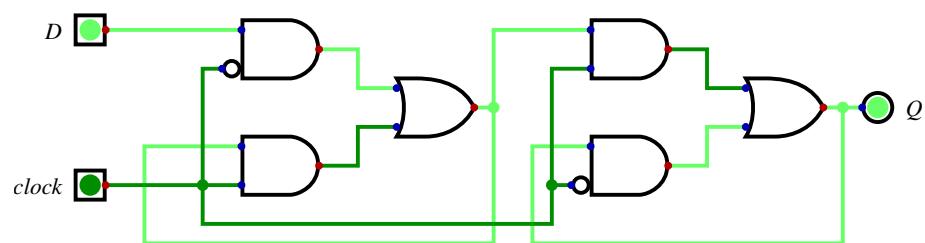
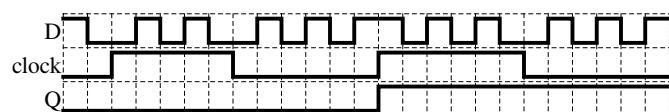


Figure 24: edge triggered D flip-flop

of them will be sampled at the same time on the rising edge of the clock. The only remaining worry related to time is to check that these edges are sufficiently far apart that the combinational circuit has time to finish its job. This is why we say that one processor can run at 1.5 GHz while another can go up to 3.2 GHz.

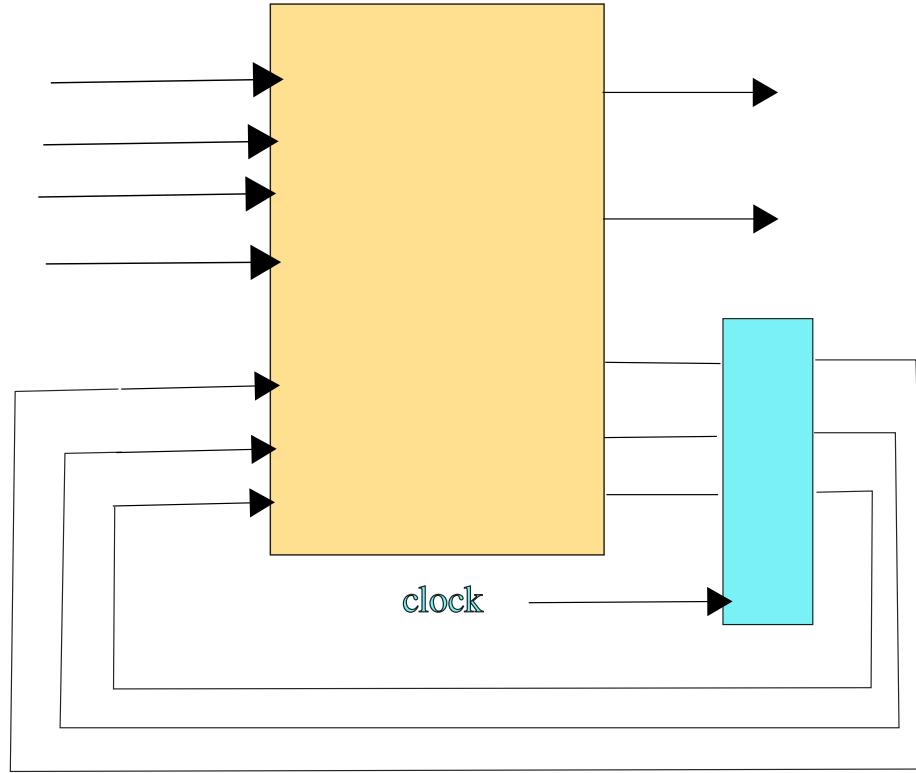


Figure 25: finite state machine

With the introduction of a lock, the state of the system can be thought of as jumping instantly from one value to the next. The number of possible states is limited by the number of bits used to represent the current state, but the number of actual states can be much smaller than that. We call this kind of system “Finite State Machine” (FSM).

One possible representation for FSMs is a table with, for example, one line for each state and one column for each possible input. Each cell would indicate the output as well as the next state. *Digital* has an editor for a very popular graphical representation: each state is shown as a circle and arrows connecting the circles show the possible transitions. The arrows indicate what values the inputs must have to jump to the other state as well as what the values of the outputs should be.

Among the examples supplied with *Digital* we have *rotDecoderMealy fsm* with *A*

and B as inputs and L and R as outputs. Each state has a name, but it also has a number from 0 to 6 which are the values that the registers should have for each state. This system can detect if an axis is rotating to the left or to the right.

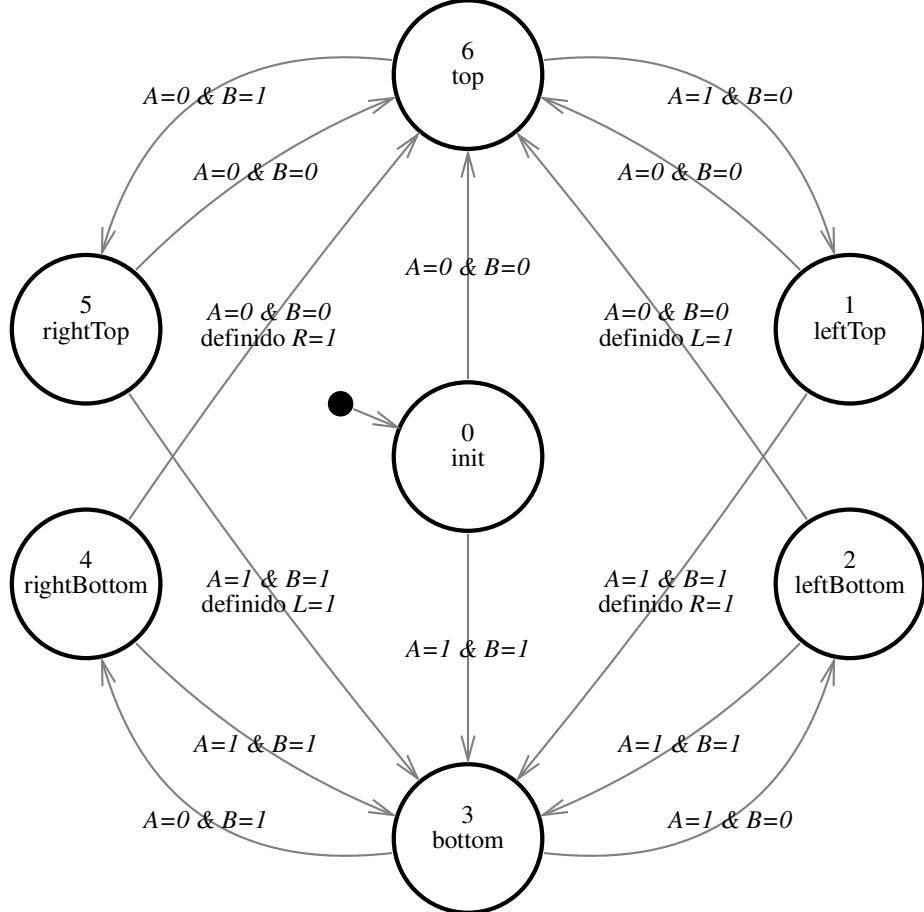


Figure 26: Mealy FSM example - rotational decoder

The “Mealy” in the file name is to indicate that this a Mealy type FSM, which was defined by George H. Mealy in 1955. One feature of this kind of FSM is that the outputs depend not only on the current state but also on the other inputs. This means that the outputs might change between clock edges if the inputs change.

In the FSM editor we have the option of generating the truth table (for the “transitions”, since as a sequential system the machine as a whole can’t be described by a truth table) and from the table we can ask for a circuit to be created. The current state is stored in registers $2n$, $1n$ and $0n$ and we can see the A and B inputs we well as the combinational circuit in the form of sum of

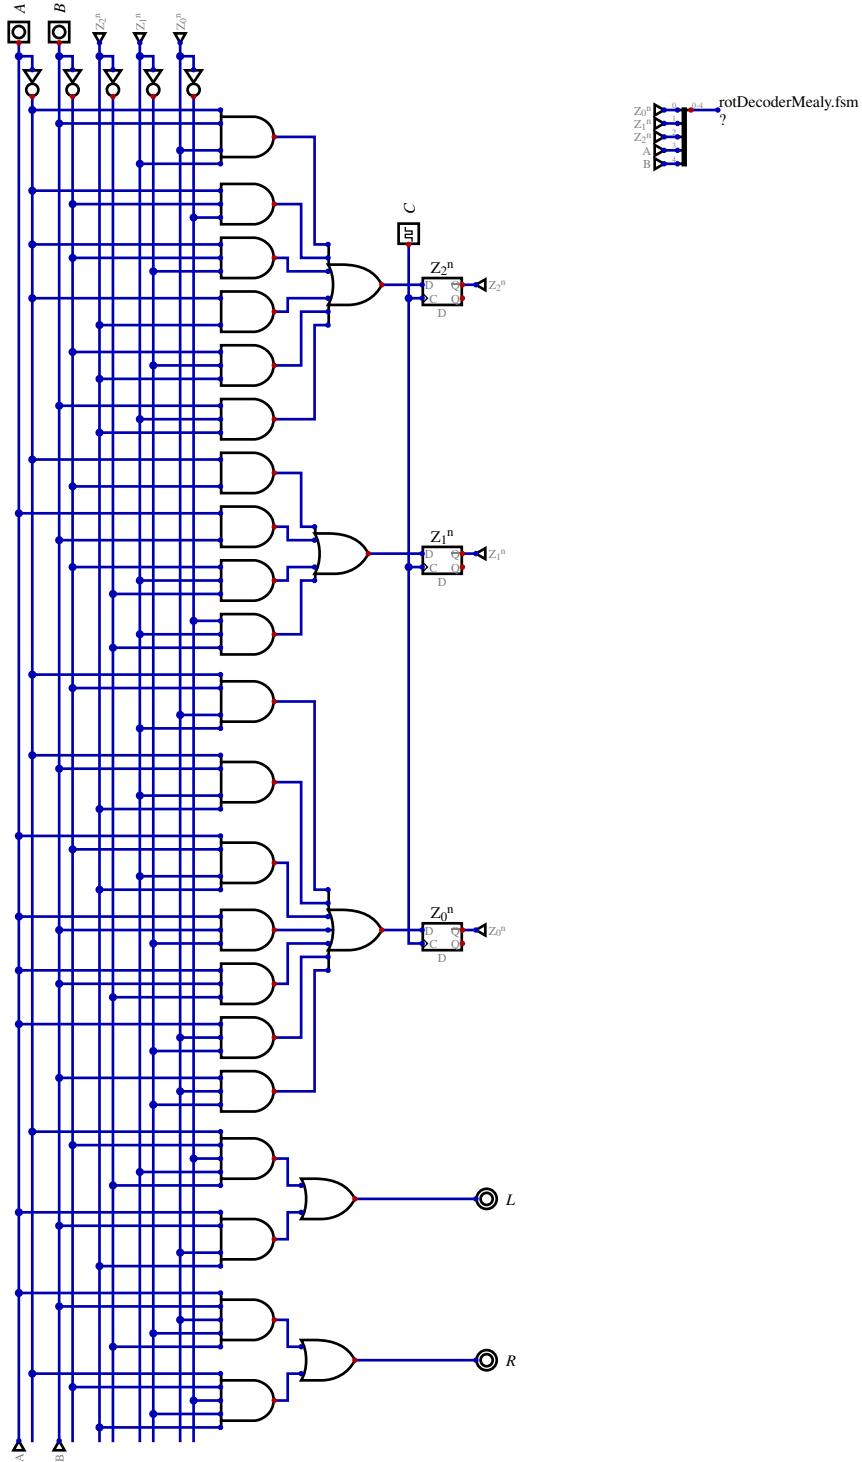


Figure 27: Mealy FSM example - rotational decoder
28

products that generated both the value of the next state and the outputs L and R .

The outputs of the registers should have wires going to the inputs of the combinational circuit. *Digital* has a feature called *tunnel* which allows a signal to jump from one place to another (or several others) without crossing over the rest of the drawing, and that was used here.

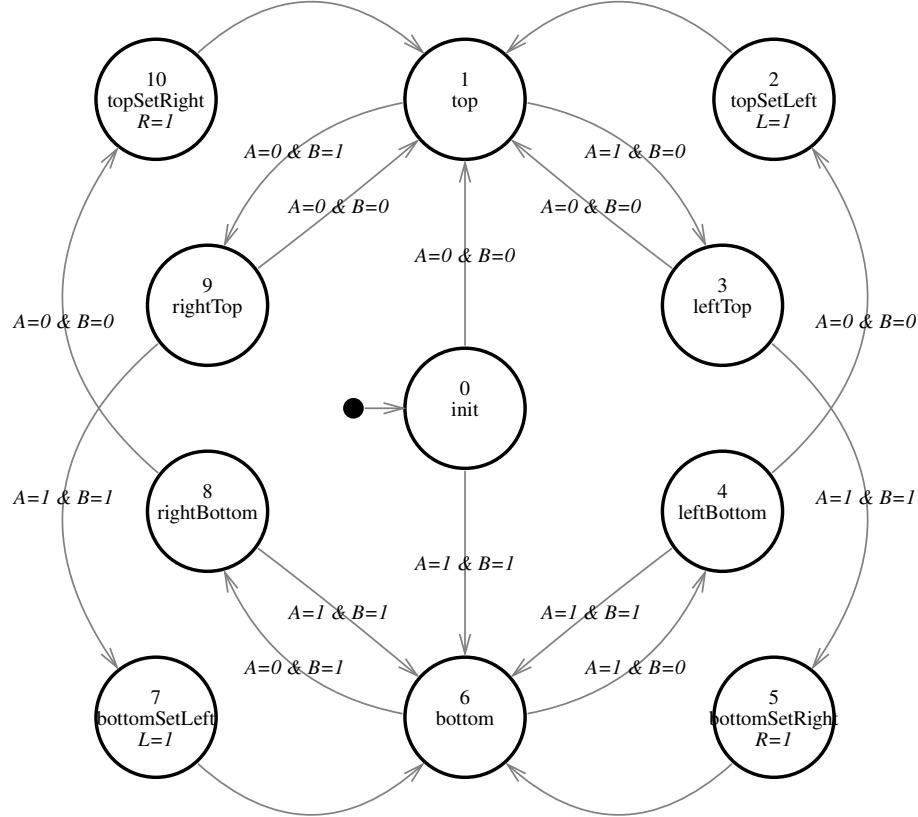


Figure 28: Moore FSM example - rotational decoder

In 1956 Edward F. Moore defined another variation of the FSM where the outputs depend exclusively on the current state. The advantage is that the outputs remain stable between clock edges. In the same example implemented as a Moore type FSM we see that extra states are needed for the same result, in this case.

The most obvious difference between this circuit and the previous one is that outputs L and R don't have any connections with inputs A and B .

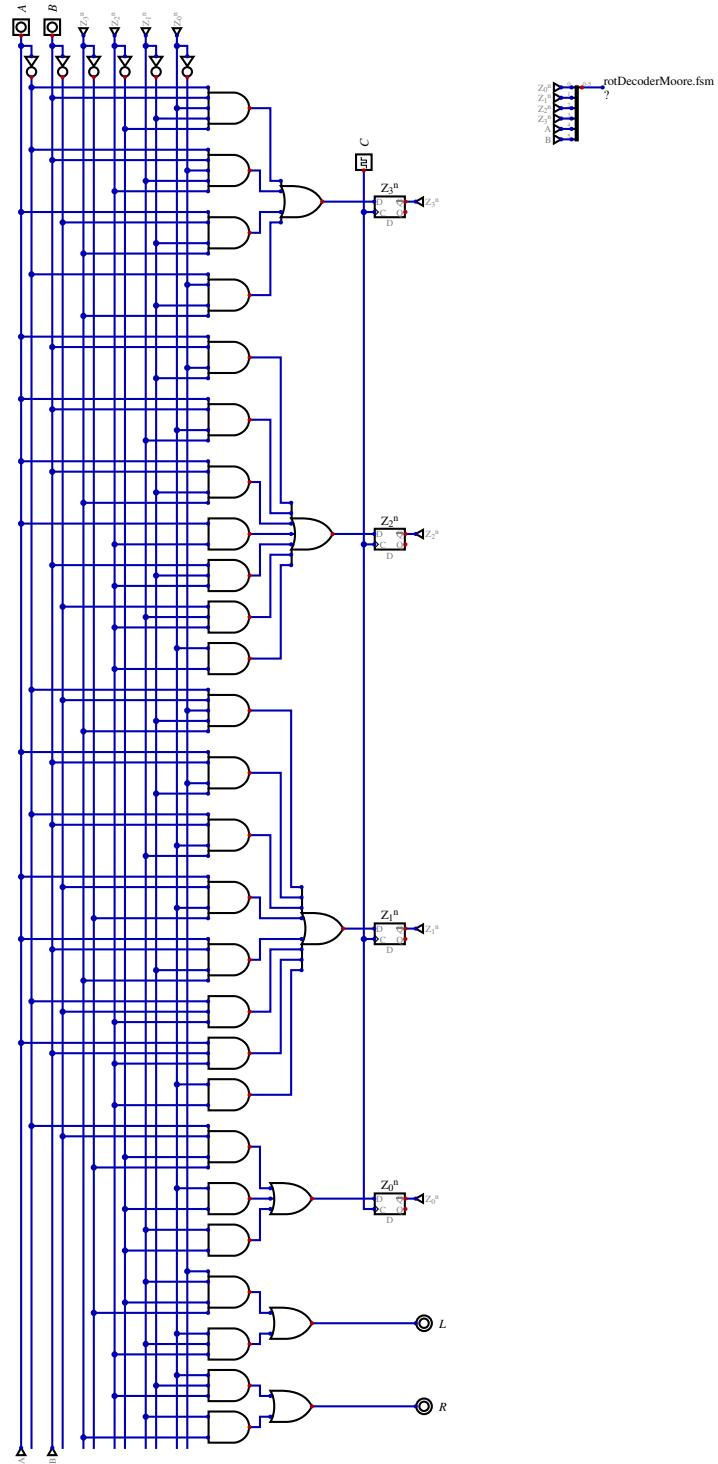


Figure 29: Moore FSM example - rotational decoder
30

3. Processors

All computational problems can be solved with a FSM. In theory. In practice even relatively small problems might need an absurd number of states and the corresponding machine would be impossibly expensive to build (if not absolutely impossible, for example if it needs more components than there are atoms in the universe).

Imagine a FSM that receives six characters of 7 bits each and which is supposed to print these characters in the reverse order. It will have 8_865_353_597_185 (8 trillion) states. The problem is that the system's only memory is the register for the current state and using that to store what characters have already been seen is not efficient. $6 \times 7 = 42$ bits to store the characters while 8 trillion states need 44 bits to represent them.

If we use a FSM connected to a small external memory it would be much more reasonable. A memory with 8 words of 8 bits each would be more than enough and a FSM with 12 state would be sufficient to control it to solve the problem.

In his 1936 paper, Alan Turing imagined something even simpler than a memory: he connected his FSM (represented in the paper as a table) to an infinite tape with individual cells that can hold a single symbol chosen from some alphabet. There is a read/write head that is positioned on one of the tape's cells. The input to the FSM is the symbol in the current cell and the outputs are a symbol to be written (possibly the same one or we don't wish to change the tape at this time) and optionally a command to move the head to the cell on the left or on the right.

Today this is known as a “Turing Machine”. A very interesting simulator which is available on the web is (<https://turingmachine.io/>) which includes several examples.

Here we have the multiplication of two binary numbers. The same thing as a pure FSM would have a much larger number of states, while here only 21 are needed. But this was created only to study a mathematical problem of the kind “there exists a turing machine capable of...” and not to be something practical. Each problem requires the construction of a new Turing Machine, but towards the end of the paper there is a very interesting proposal: a Universal Turing machine that would receive on the same tape as the input data a representation, in the form of a sequence of symbols, of a Turing Machine that would solve the selected problem. We now call this an “interpreter”. Once such a machine was built, changing the tape would modify its behavior. This is what we call “software”.

von Neumann and Harvard Architectures

The first computer (or “electronic brain” as the press initially called them) was ENIAC in 1946 (previous projects were kept secret for many years). Designed

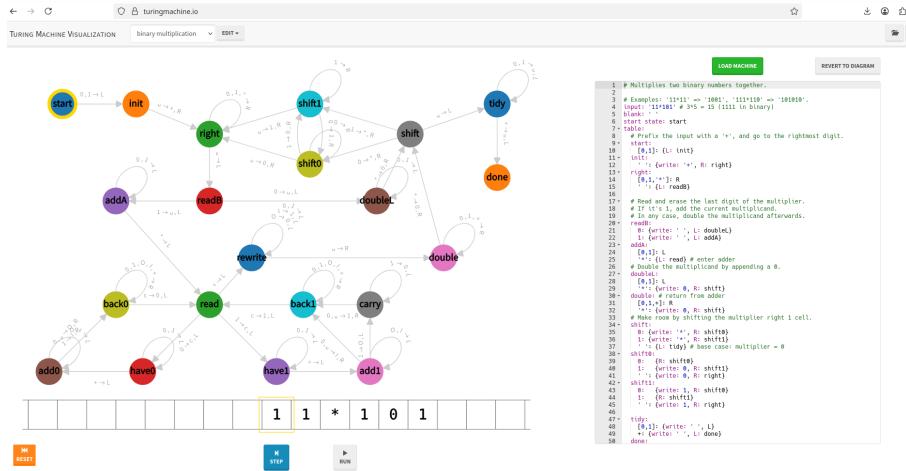


Figure 30: Turing Machine simulator

by John Mauchly e J. Presper Eckert at the University of Pennsylvania, a key limitation of the ENIAC was the need to reconfigure the hardware using patch cable panels for each new problem. So even during its development they started discussing its successor, to be called EDVAC.

One of the participants of these debates was John von Neumann and he wrote a detailed report on these ideas. Another participant, Herman Goldstine, ended up distributing the report to several external grupos with von Neumann as the sole author, so this style of computing is known as the “von Neumann architecture” even though it was created by a group.

John von Neumann liked analogies with biology and so called the parts of the computers “organs” and where the data is kept “memory” (others, specially IBM, preferred terms such as “storage” but ended up losing this battle).

The control unit is just a FSM, which we have already seen. The arithmetic and logic unit is a more complicated version of the adder/subtractor which we have also already seen. The input and output are different for each computer, so for now we will ignore them.

The memory stores both data and programs (just like the tape of the Universal Turing Machine). One bit of this memory is like a register, which we have already seen. What we didn't talk about is how to select one bit from many, but the multiplexer is similar to the mechanism used. An interesting feature of the von Neumann architecture is that the central processor (CPU) can be made from a technology that is completely different from that used to make memory. Some examples:

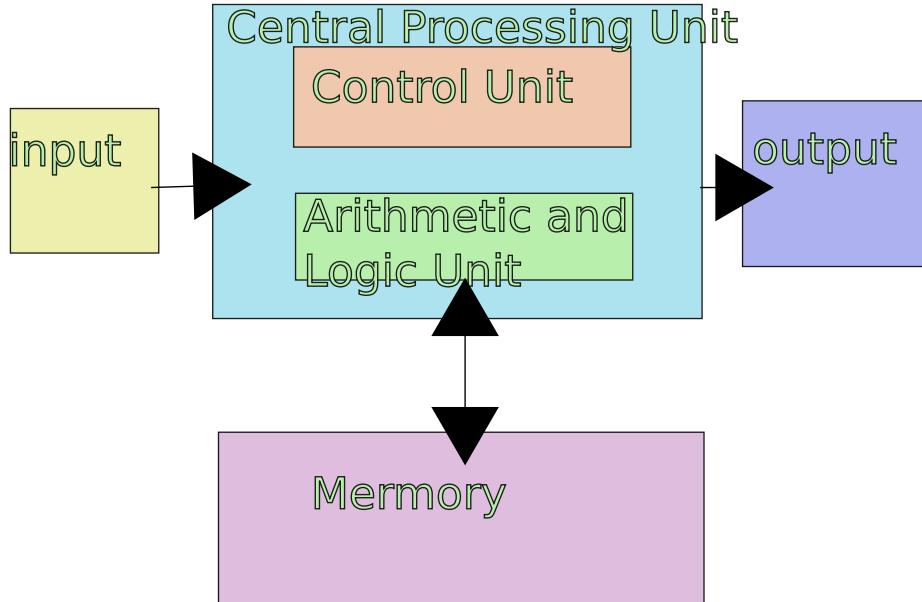


Figure 31: von Neumann architecture

Computer	CPU	Memory
EDSAC	vacuum tubes	mercury tanks
IAS	vacuum tubes	Williams tubes
LGP-30	vacuum tubes	magnetic drum
PDP-8	transistors	magnetic cores
modern PC	digital chips	chips with vertical capacitors

IAS is “Institute for Advanced Studies” in Princeton where John von Neumann and his team built the computer from his report. That is why it is more rarely referred to as the “Princeton architecture”. That is to contrast it with the “Harvard architecture”, which is named after a computer that IBM built for Howard Aiken at Harvard University. The only difference is that the data memory and program memory are separate. The separation allows an instruction and some data to be read at the same time, which can simplify the project. But makes it impossible for a program to modify itself.

Today all computers except for the simplest ones are hybrids: directly connected to the processor we have two tiny memories known as the level 1 instruction cache and the level 1 data cache. These memories only store copies of information that have been recently used by the computer. When the needed information is not there, a level 2 unified cache is accessed. And there might be a level 3 cache and then finally the main memory as in von Neumann’s drawing. This combines

the hardware advantages of the Harvard architecture with the programming advantages of the von Neumann architecture.

We mentioned that inputs and outputs are specific to each computer. In early computers this was reflected in the instruction set. A computer might have one instruction to read a key and another instruction to write to the tape. Around 1970 the idea of making input/output devices appear as special memory locations began to become popular. This allows the processor to use “normal” instructions for everything. Of the processors current used, only x86 (Intel and AMD) still have special instructions for input and output and all the users use “memory mapped peripherals”, which we will do as well.

MCPU16h

In the EDAVC report the idea was put forth of representing a program as a series numbers where the most significant bits would represent an “order” being given to the computer (this is now called the “operation code” or just “opcode) and the least significant bits would be the address of the memory position to be used for this instruction. Many operations need two operands and produce a result that must be stored somewhere. This would require 3 addresses, mas EDVAC inherited from mechanical calculators the idea of a special register called “accumulator” which supplies one of the operands and is the destination of the result for most of the instructions.

How many different instructions will we have? This defines how many bits we will need for the “opcode”. It turns out that it is actually possible to do everything with just a single instruction, but that is not a good idea from an educational viewpoint. An example of such an instruction is SUBLEQ which has 3 addresses and subtracts the value at the first address from what is at the second address (saving the result back there) and if the result was negative or equal it jumps to the third address. Programs written for such a computer are nearly as hard to understand as those for a Turing Machine.

A more reasonable alternative is MCPU with its 4 instructions (two bits for the “opcode”). Tim Böscke was inspired by MPROZ and its 3 instructions, but replacing the memory to memory operations with an accumulator in the EDVAC style greatly simplified the project.

MCPU16h has two differences: while the original MCPU is 8 bits wide so only 6 bits are left for the address (pny 64 bytes, which is sufficient for the simplest examples) MCPU16h, as its name implies, is 16 bits wide and its 14 bit addresses can handle 16 thousand words of 16 bits each (32KB). And the “h” at the end of the name is from Harvard. But like the EDVAC, MCPU16h depends on being able to modify a program while it is running but that is not possible with a Harvard architecture. But *Digital* has a two port memory and we will use that to make it look like there are two separate memories but what is written to one can be read by the other.

The choice of the Harvard architecture was to allow the use of a combinational circuit for the control unit. In a von Neumann architecture (like the original MCPU) reading an instruction happens in one cycle while reading data in a different one. Because of this the control unit has to be a sequential circuit, which is a bit more complex.

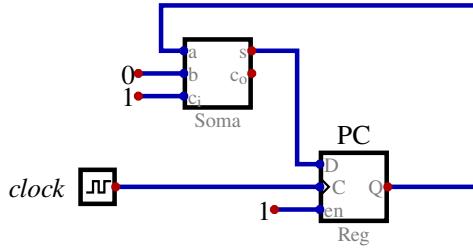


Figure 32: PC do MCPU

With this simple circuit, at each rising clock edge the PC advances to the next word. In this case we connected the “enable” of the PC register to 1 since we never stop changing the PC. But in a more complete project there will be situations where that will happen (if we have to wait for the instruction memory, for example).

So let’s add the dual port memory. Port 2 is limited to reading and will be used as the instruction memory. The control signals for port 1 are connected to constant values that won’t interfere with the operation (not writing anything, for example).

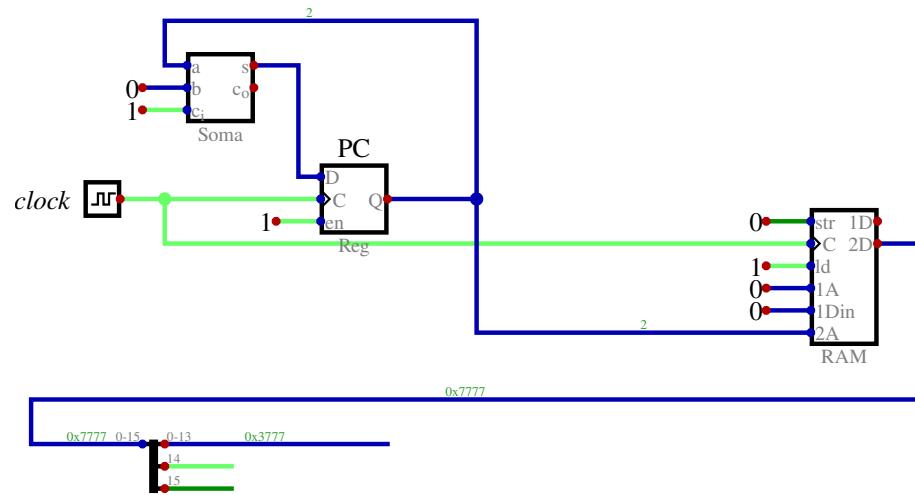


Figure 33: dual port memory

When simulation starts the memory is full of 0s. Later on we will use the contents

from a file to fill in the memory, but for now we are manually editing the values of the first few words every time the simulation starts just to see that something is actually happening. The picture shows what happens after two clock rising edges.

The data arriving from memory (0x7777 indicating that it is the hexadecimal equivalent to the binary 0111011101110111) is divided into a 14 bit address (0x3777) and two “opcode” bits (0 and 1). We are decoding the instructions using only wires.

We shall implement the first of the four instructions: JCC (“Jumps if Carry Clear”). We need to choose an “opcode” for this instruction and we selected 1 and 1.

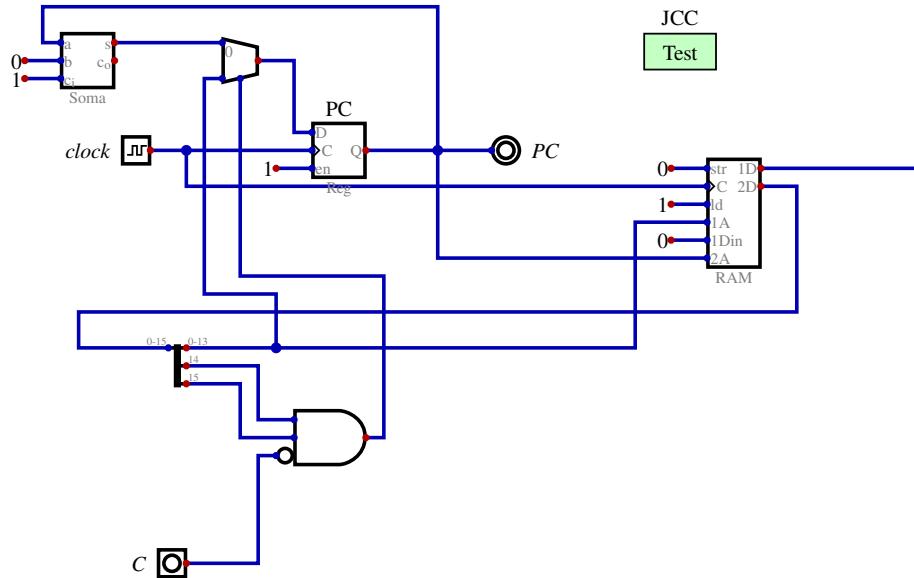


Figure 34: JCC instruction

The multiplexer decides between jumping to the address indicated by the instruction or just adding 1 to the current PC. The first case should only happen if the opcode is that of the JCC and if the C bit is 0. The 3 input AND gate detects this situation and controls the multiplexer.

Here we are creating a single circuit, but normally the AND gate would be part of the control unit while the multiplexer, the PC register and the adder would be part of what we call the “data path”. And the memory would be separate from the CPU.

```
clock C PC
program(0x0000, 0xC006, 0xC001, 0xC008)
0 1 0
```

```

C 1 1
C 1 2
C 0 1
C 1 2
C 0 1
C 0 6

```

To test this circuit we need to make PC an output so it can be compared during the test. We initialize the first 4 words of memory with a tiny test program. Instruction 0 will be ignored for now and instructions 1, 2 and 3 are all JCC. The first time that instruction 1 is executed (line 5 of the test) $C = 1$ and PC is incremented to 2. In lines 7 and 9 instruction 1 is executed again with C equal to 1 and C equal to 0 respectively. Only in the last case it doesn't go on to 2, but instead jumps to 6.

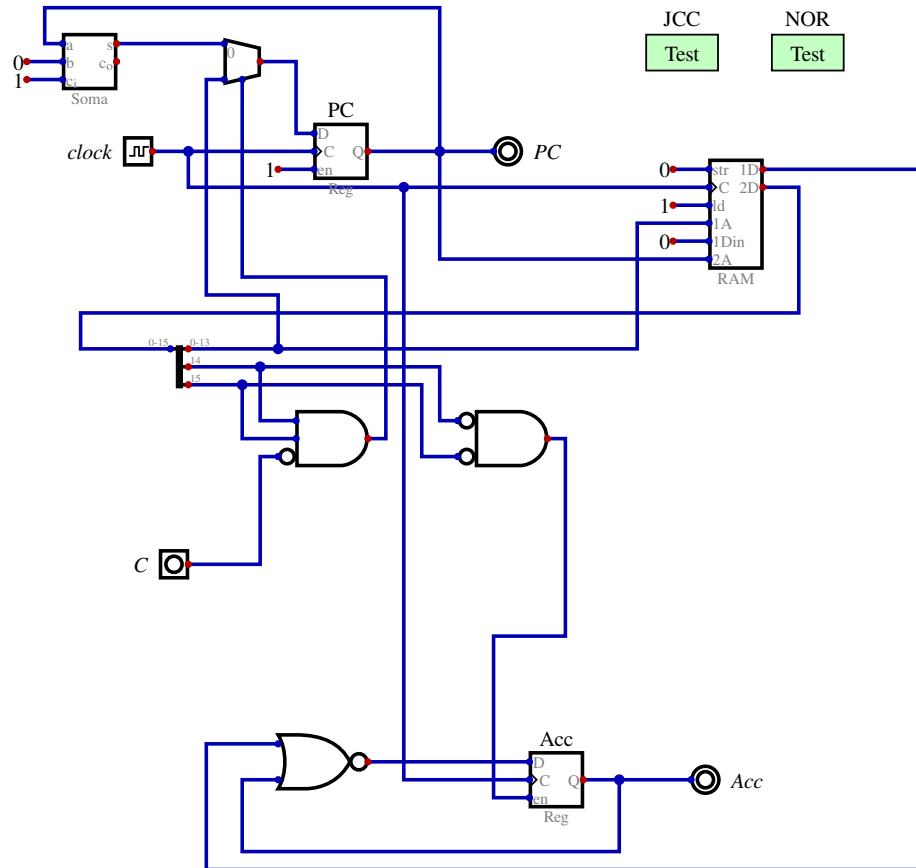


Figure 35: JCC and NOR instructions

The next instruction we will implement is NOR, for which we chose the opcode

0 and 0. We need a new 16 bit register that we will call Acc. We will also have an output with this name to be able to use it in the tests. The 16 bit wide NOR gets one input from Acc and the other from memory, with the result going back to Acc.

First we test JCC to check that it continues to work. Then we create a new test for NOR. The circuit must pass both.

```
clock Acc
program(0x0004, 0x0000, 0x0001, 0x0002, 0xFFFF)
O 0
C 0
C 0xFFFFB
C 4
C 0xFFFFA
```

Curiously, words 0, 1 and 2 are initially used as instructions but then are used as data.

The STA instruction (“STore Accumulator”) will have the opcode of 1 and 0. Acc’s output will be the data input for the memory (which was always 0 up to now). And the memory’s *str* signal will be activated by this instruction, writing to the indicated address.

After checking that the JCC and NOR tests continue to work, we create a new test for STA.

```
clock Acc
program(0x0004, 0x8004, 0xC000, 0x0000, 0xFFFF)
O 0
C 0
C 0
C 0
C 0xFFFF
C 0xFFFF
C 0xFFFF
C 0
C 0
C 0
C 0xFFFF
```

It isn’t possible to test STA without also using other instructions. The results of writing data is only visible when you read back that data and for that we need NOR. The test also uses JCC just to keep the code short - a sequence of NOR, STA, NOR, STA, NOR... would be good enough test.

The last instruction, ADD, would be nearly the same as NOR if it wasn’t the complication of the carry bit. But as both ADD and NOR write to Acc we need a multiplexer to decide between them.

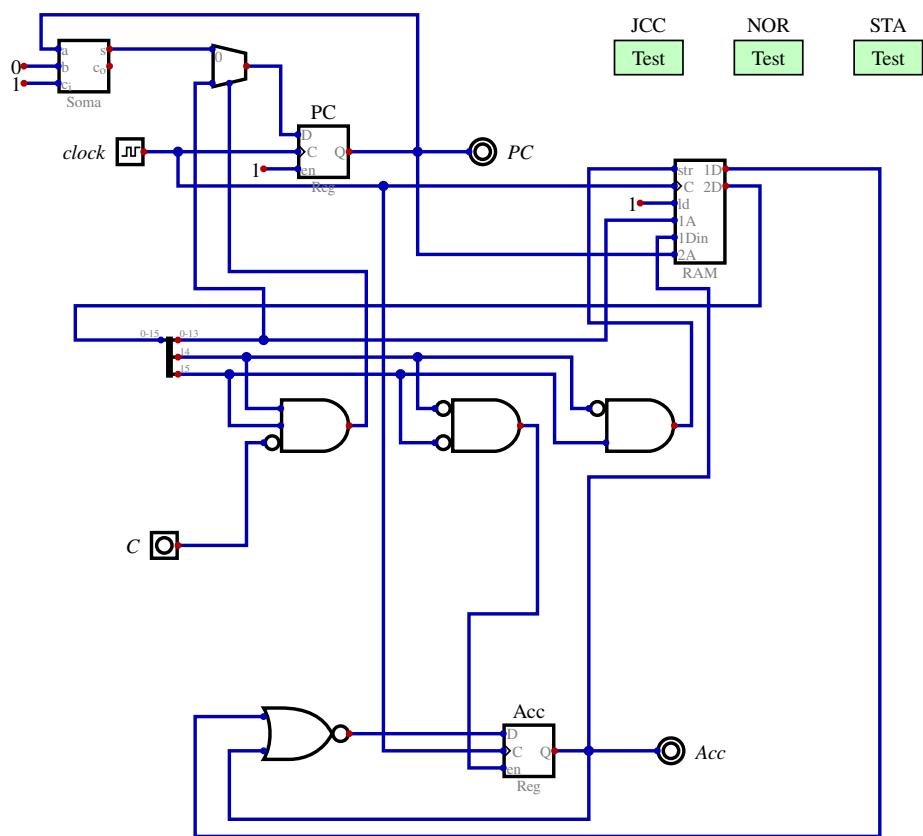


Figure 36: JCC, NOR and STA instructions

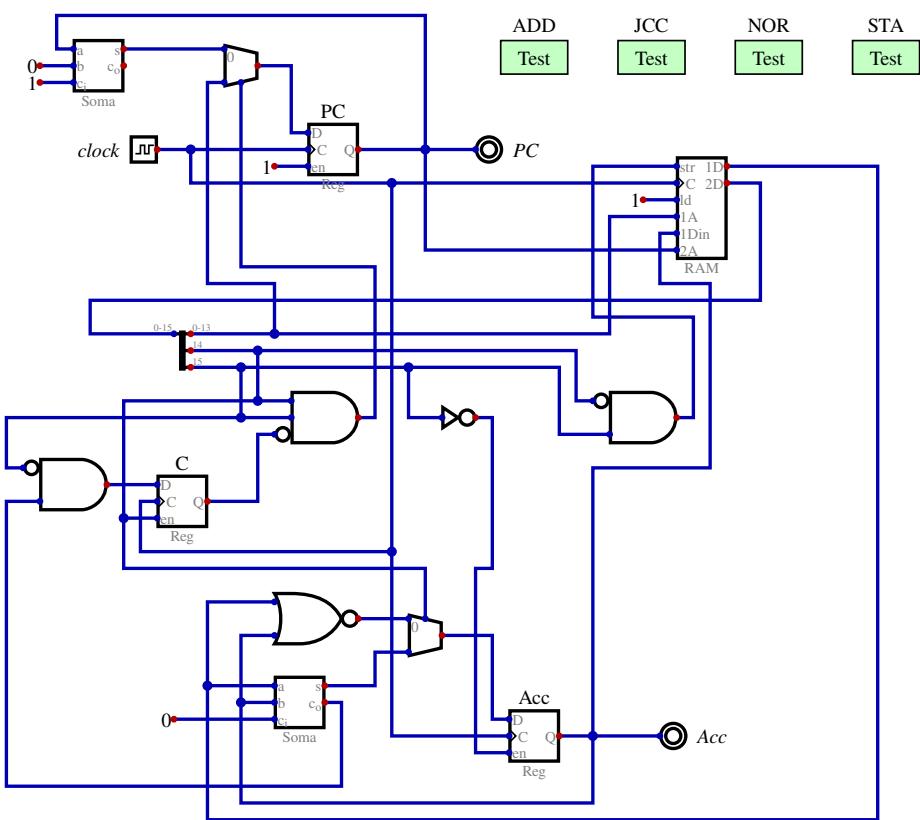


Figure 37: the complete mcpu

Now Acc needs to be enabled both for ADD (opcode 01) as well as NOR (opcode 00), so we can replace the AND gate with the two inverted inputs with a simple inverter. The new adder has the same inputs as the NOR.

The C input is replaced by a 1 bit wide C register. This forces us to change the JCC test. The definition of JCC is that C is cleared independently of whether the jump happens or not. This is used so two JCC in a row will be an unconditional jump (which in other processors is a separate instruction). The carry output of the adder is forced to 0 in the case of a JCC instruction (or a STA, but in that case it doesn't matter).

C needs to be enabled for the ADD (opcode 01) and JCC (opcode 11) instructions, so the lower bit of the opcode can be used directly for that.

With that, MCPU is complete and can execute complex programs.

Software

We have already written small programs for the MCPU16h in the tests we have created. But indicating each instruction directly as a hexadecimal or binary numbers, what we call “machine language”, is unfeasible for programs larger than 10 instructions or so. Fortunately, a program called “assembler” can read a text file where each line corresponds to one instruction and some letters indicate the operation (“NOR”, “ADD”, “STA” and “JCC” in the case of MCPU16h) and generate the corresponding machine language. We can mark certain lines using textual labels and then use the same text as the address field of other instructions so the assembler will take care of the calculation of the actual value associated with each label.

Fancier assemblers, like *GNU as*, allow the definition of “macros” which are texts (possibly with some arguments) which are then expanded everywhere they are used in a program. In this project we use this feature to get one version of *as* (for the x86 processor, for example) to generate machine code for the MCPU16h. These macros are found in the file *mcpu16.inc*, which has near its start:

```
absStart:
    .macro NOR a
    .word (0x3FFF & ((\a-absStart)/2))
    .endm

    .macro ADD a
    .word (0x3FFF & ((\a-absStart)/2)) | 0x4000
    .endm

    .macro STA a
    .word (0x3FFF & ((\a-absStart)/2)) | 0x8000
    .endm
```

```
.macro JCC a
.word (0x3FFF & ((\a-absStart)/2)) | 0xC000
.endm
```

A program that includes this file can have something like “JCC myLoop” and the corresponding 16 bits will be generated. One complication is that *as* considers a label like “myLoop” to have relative value and can’t, in principle, calculate the expression which generates the instruction. The idea is that a program might be organized as several modules and the assembler translates one at a time and a later step links all the modules into the final program. This means that the addresses will only be known after that. But the programs we will write for MCPU16h will be relatively simple and we will not have more than one module. This allows us to define “absStart” as the first thing in the program so the expresison “myLoop-absStart” can have a value that is known to the assembler.

Besides the 4 instructions that the hardware actually understands, we can use macros to define new instructions as short sequences of these instructions. For example: CLR (clears the accumulator to zero), LDA (loads the accumulator with a value), LDP (loads a pointer into the accumulator), LDI (loads into the accumulator indirectly), NOT (inverts the accumulator), JMP (jumps unconditionally), JCS (jumps if *C* is 1), SUB (subtract the accumulator from a value) and CALL (calls a subroutine). The use of these “pseudo-instructions” generates programmas that are somewhat large. Programas for MCPU16h are around 10 times larger than for more reasonable processors.

The pseudo-instructions IN and OUT read and write to the last address in memory and it is up to the circuit to detect this and read and write from a peripheral instead of memory. Two additional macros, STARTCOUNT and COUNT, were defined to control a circuit used to measure speed so we can compare different processors. In this project we will not implement this circuit so these pseudo-instructions will just access the next to the last word in memory without any side effects.

The command “make ALLHEX” generates the .hex files we will use to initialize memory in our simulations. These are saved in the *hex* directory using the .S source files in the *soft* directory. Programs *as* and *objcopy* are used by *make*. Later on we will use *riscv32-unknown-linux-gnu-as* instead of *as* to handle programs for RISC-V. We use .s (lower case) in the source file names in this case so *make* can decide which assembler to use.

4. FPGAs and Shin JAMMA

5. Video and Audio

6. Pegasus 42

A. History

From 1977 to 1991 Brasil had a “reserved market policy” for small scale computers. At the time nearly all imports were forbidden, from automobiles to chocolate including computers. But for other products a foreign company like Volkswagen or Nestlé could open factories in the country and sell locally. The reserved market didn’t allow this for the case of mini and microcomputers.

In times of high inflation and no economic growth, the computer sector was one of the few that was greatly expanding. This attracted investment from non related areas, specially due to the opportunity of copying foreign products without having to deal with competing with the originals. It was in this context that a group from the financial sector contacted the lawyer Amaro Moraes e Silva Neto for ideas, and Amaro told them that one thing nobody was making was a children’s computer. He said he knew a computer genius who could develop such a product.



Figure 38: Amaro Moraes e Silva Neto

He was thinking of Fábio Cavalcante da Cunha, who was officially an electronic engineering student at the Polytechnical School of the University of São Paulo (Poli-USP) but was then more focused on his job at a computer store.

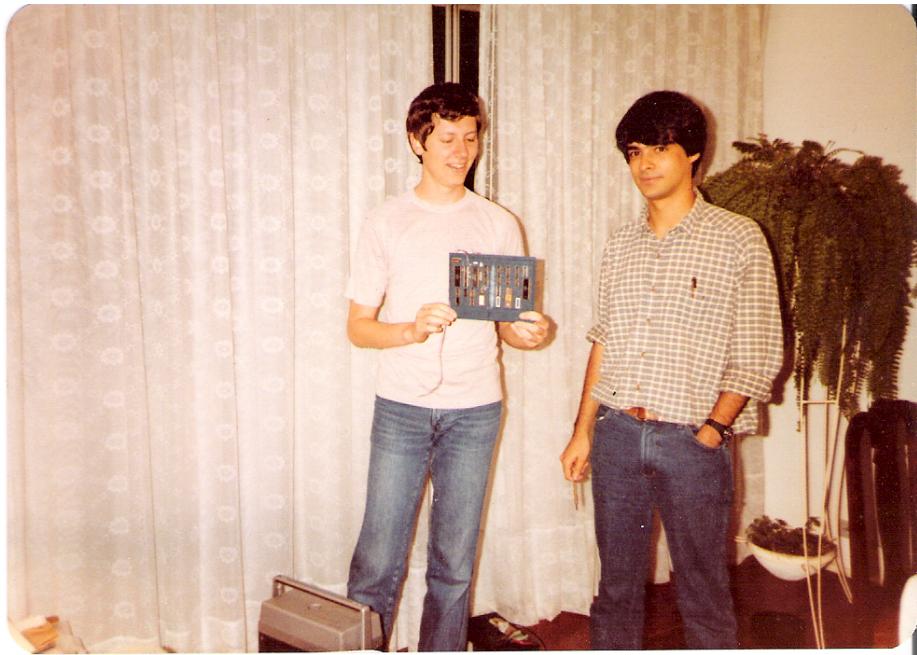


Figure 39: Jecel and Fábio with the prototype's front side

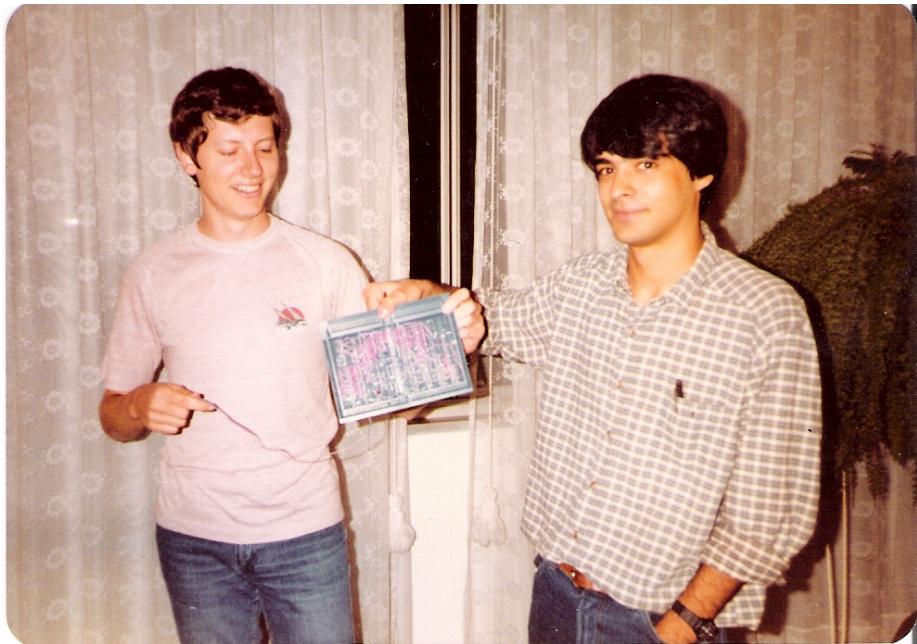


Figure 40: Jecel and Fábio with the prototype's wiring

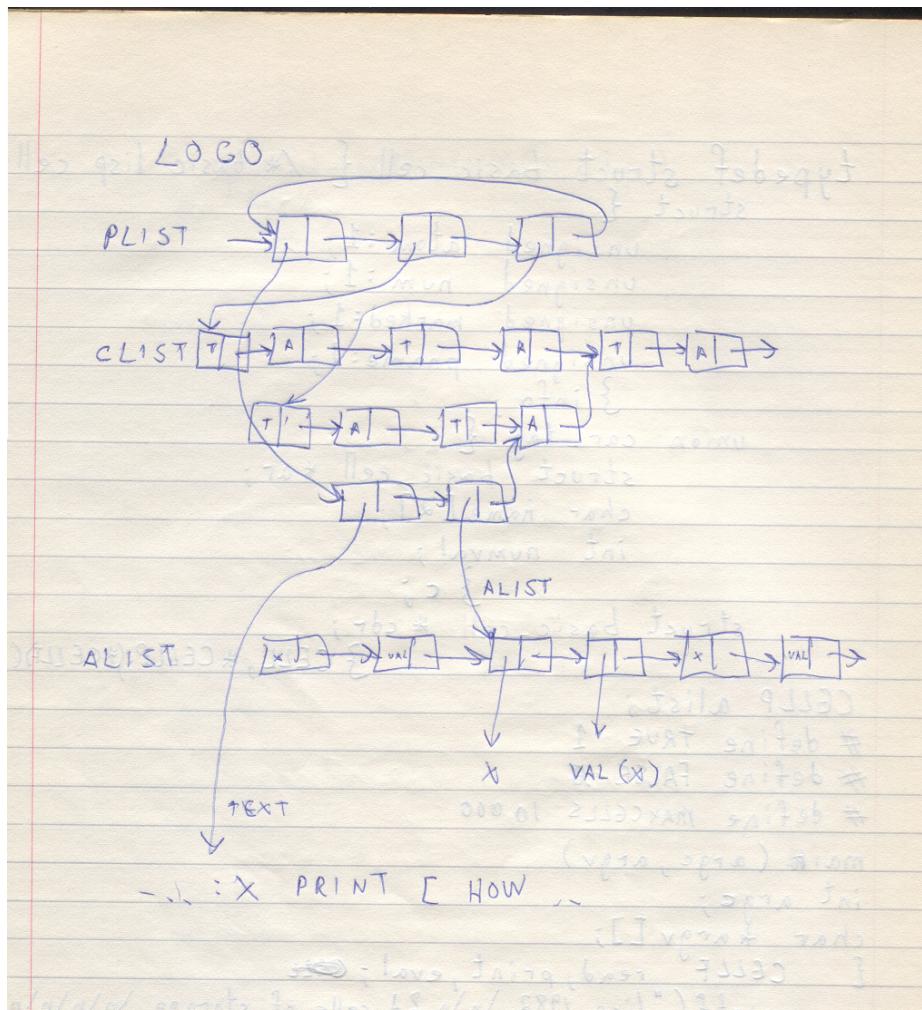


Figure 41: Logo's structure

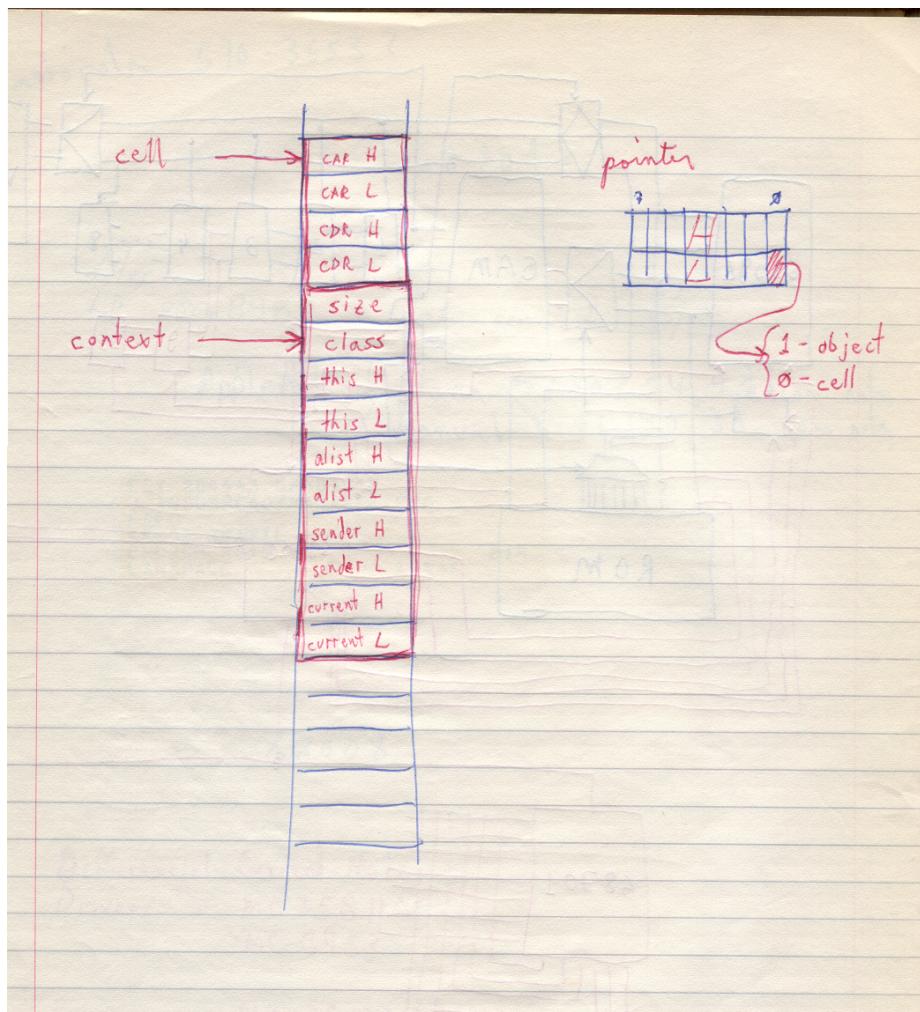


Figure 42: object format

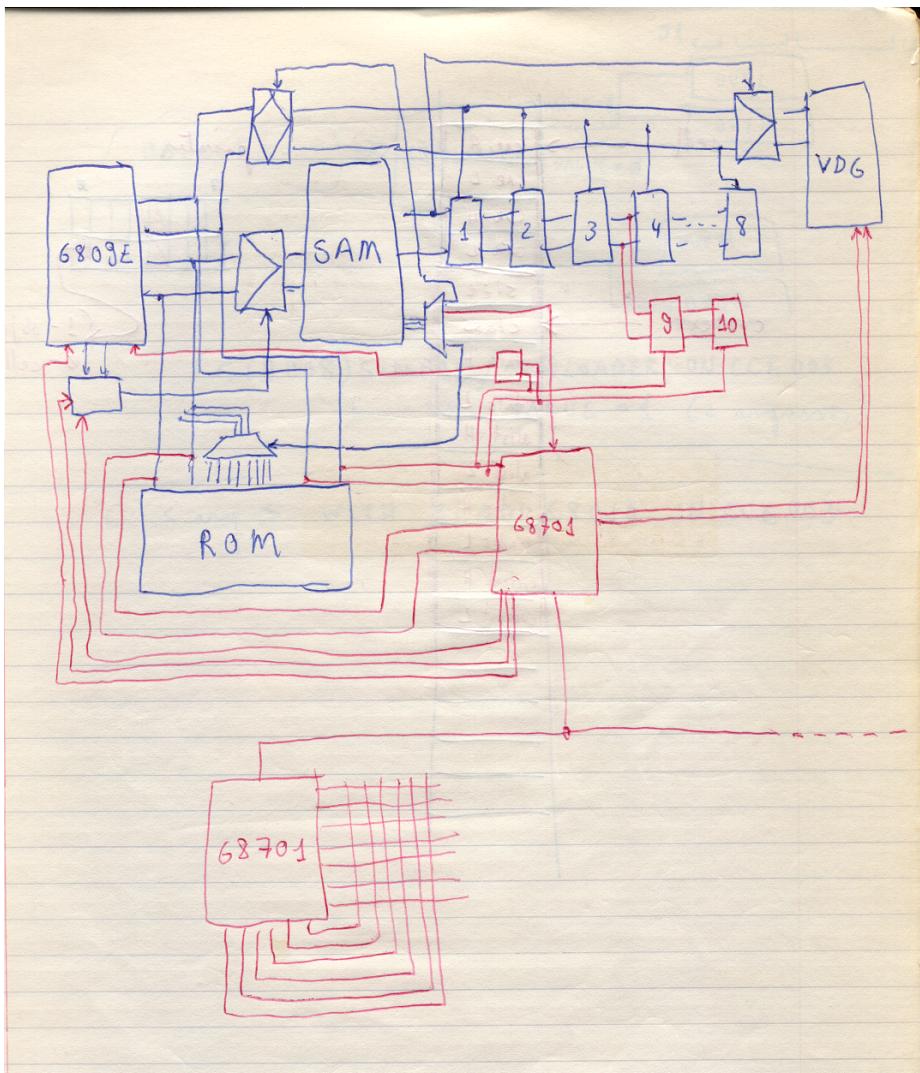
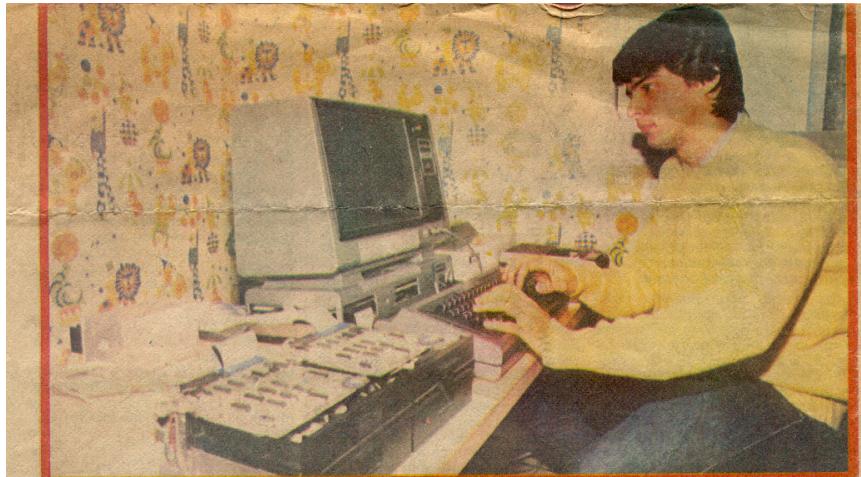


Figure 43: Pegasus schematic



Com linguagem Logo, a máquina de Fábio é inédita

19 anos, aficionado, Fábio inventa a nova eletrônica

Com seu computador, ele se torna um pioneiro

JOSÉ EDUARDO MENDONÇA

Influenciado pelo pai, um rádiodoador, Fábio Cavalcanti resolveu, aos sete anos de idade, desmontar um rádio para ver como as válvulas funcionavam. Ganhou uma bronca, e o pai fez um rádio que nunca mais funcionaria. Além disso descobriu, com a experiência, a paixão pela eletrônica que o levou agora, aos 19 anos, a fabricar seu próprio computador.

Aos 13 anos, Fábio fez um curso de computação e o "arquivou na cabeça", devido a interesses que julgava mais apropriados no momento, como seu autorama e os aeromodelos. Mais tarde, no colégio, começou de novo a "ler revistas especializadas e montar pequenos circuitos". Tentou fazer um curso mais avançado e os que encontrou "limitavam demais o conhecimento". Mais uma vez, esqueceu temporariamente sua paixão para enfrentar o cursinho e entrar na Politécnica no curso de engenharia eletrônica, depois de já ter conseguido uma vaga na universidade de São Carlos, para onde migrou abandonando sua cidade natal, Londrina.

Seu computador, o Projeto Pegasus, para a linguagem Logo (ver página "Computadores"), começou a nascer com o exercício de programação em calculadoras e a busca constante de maiores informações. Na época, "78/79, ainda não existia este clima de hoje com os computadores, só mexia com eles o pessoal de processamento de dados em grandes computadores". Fábio aprendeu "tudo que podia saber uma calculadora, depois de fugir durante um ano", e optou pelo caminho natural: em 1983 foi para os Estados Unidos, mais especificamente para o coração da indústria de computação, o laboratório do Massachusetts Institute of Technology, na Universidade de Stanford, em pleno Vale do Silício.

Um de seus encontros lá foi tão emocionante quanto o encontro, digamos, de um guitarrista daqui com Frank Zappa. Fábio trocou horas de papo com seu guru Marvin Minski, um dos papas da inteligência artificial, e conseguiu provar para si mesmo "que podia estar lá". A mágica foi não ter conseguido uma bolsa de graduação e ter "assumido o caminho de volta".

"Se eu tinha de fazer alguma coisa aqui, senti que teria de fazer como sempre fiz: sozinho." Armado de dezenas de livros, estimulado por dois ou três colegas de escola também aficionados pelos computadores, utilizando seu TRS-80 da Radio Shack, Fábio partiu para seu projeto. Antes, outra deceção: ofereceu seus serviços à Poli, disse que tinha estado no MIT e que gostaria de colaborar. Ouviu como resposta que suas idéias não tinham aplicação comercial imediata. "Parece que dentro da escola não acreditam muito no valor dos alunos, e assim nós temos de criar fora dela."

O computador de Fábio será, diz ele, "possivelmente o primeiro do mundo a vir já com a linguagem Logo, e não com a Basic." A idéia: "Aproximar as crianças dos computadores. Tem cor, som, pode até mesmo programar jogos e serve para crianças a partir de 4 anos até a adolescência, ou mais. Há também excelentes recursos para aplicações mais maduras." O projeto Pegasus ainda não foi enviado para a necessária avaliação e aprovação da Secretaria Especial de Informática, mas, segundo um de seus autores (o outro é o colega Jacel Mattos de Assumpção, também um autodidata, e "um dos poucos gênios" que Fábio conhece), "a máquina já está muito bem definida".

O projeto Pegasus deverá ser comercializado por cerca de 300 mil cruzeiros, tem 64 k bytes de memória, 8 bits e utiliza um chip 6809 da Motorola. Fábio enfatiza que seu computador

não é de forma nenhuma um "game", passatempo que critica: "O videogame pode estimular, mas para uma posição errada. As pessoas podem aprender a utilizar o computador como brinquedo, e apenas assim. Ai ele não estará servindo como ferramenta de aprendizagem."

Além de estudar na Poli e dar aula de computação na Emarés para crianças, Fábio Cavalcanti "troca figurinhas" com grupos de usuários de computadores Apple e TRS-80. Há, no entanto, poucos jovens. "Isto está apenas começando. Os mais novos aprendem diferente. Os mais velhos vai falar como o computador funciona, e a garotada vai perguntar como é que eu faço."

Como profissionais de outras áreas, Fábio e seus colegas têm sua galeria de ídolos e seus livros e passatempos prediletos. Exemplos de pessoas quentíssimas: Wolzniack, um dos inventores do Apple (mais "fera" que o outro, Steven Jobs, tído pelos mais ruros como comerciante), Scott Adams, genial escritor de softwares para games em computador; Leo Chrissopherson, programador de animações e Nolan Bushnell, inventor da Atari.

Todos leem Sherlock Holmes por seu gênio de dedução, os livros de Tolkien por sua magia, o já clássico "Godel, Escher e Bach", de Douglas Hofstadter, um estudo comparativo entre o matemático, o arquiteto e ilustrador e o músico. Adoram assistir "Star Trek" (Jornada nas Estrelas) e cultuam em forma de livros ou filmes a série "The Twilight Zone" (Além da Imaginação). Jogam horas seguidas o jogo de computador chamado "Adventure", tentando sempre superar suas próprias marcas. Com uma cultura semelhante a esta em todo o mundo, estes jovens aficionados dos micros e da robótica começam a criar a tecnologia do futuro.

Figure 44: news item in Folha de São Paulo
48

LINGUAGEM LOGO

Dois estudantes da Escola Politécnica da USP, Fábio da Cunha e Jecel Mattos Jr., desenvolveram e estarão lançando na Feira de Informática o PÉGASSOS, primeiro microcomputador que vem com linguagem LOGO. Segundo Fábio da Cunha, a linguagem LOGO é a mais adequada para se aprender computação por sua facilidade de uso, daí sua grande aplicação junto a crianças. O PÉGASSOS funciona com microprocessador 6809, da Eletrola, com velocidade de 1MHz. A memória do sistema é de 16Kbytes de ROM, podendo chegar até 64K, e a memória do usuário é de 64Kbytes. O novo micro tem teclado alfanumérico, com todos os elementos do português, como ç e todos os acentos. O equipamento já vem com interface embutida para ligação com gravador cassete e com televisor comum, inclusive TV a cores (8 cores no sistema Pali/M). O PEGASSOS possui sistema operacional gráfico, som, e tem expansões para disco, impressora e para ligação a um outro micro. Pode também receber cartuchos com jogos e com outras linguagens.

Figure 45: news item in Microsistemas magazine
49

Brasil coleciona casos de sucessos e fracassos

A saga de Steven Jobs e Stephen Wozniack não contagiou apenas jovens americanos. A idéia de confrontar uma grande companhia, tomando como ariste apenas o cérebro, um punhado de chips e um micro doméstico, ganhou entusiastas adeptos no Brasil, e aqui, como nos EUA, os herdeiros de Jobs e Wozniack podem ser divididos entre vitoriosos e fracassados: Ivan Nazarenko, 21, por exemplo, que estará concluindo o curso no Instituto Tecnológico da Aeronáutica (ITA) em 1986, pode ser perfeitamente enquadrado no primeiro caso. Em 1983, quando os micros fabricados no país ainda não tinham caracteres específicos da língua portuguesa, como o cedilha e acentos, Nazarenko, com apenas 18 anos, criou uma placa que "ensinava português" aos computadores nacionais, batizada de Ivana.

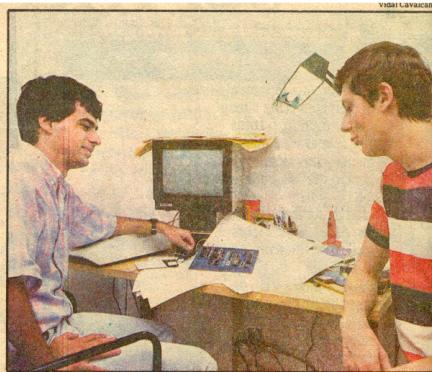
"Foi um invento simples, mas que não tinha precedentes no mundo", explica Nazarenko, que já vendeu através de sua microempresa 1.700 placas, que hoje são comercializadas a 17 ORTNS. "Não deu para ganhar muito dinheiro, mas alimentou certos luxos que tenho. Sem a Ivana, eu não poderia ter um Passat 1.8, que adoro, e seria obrigado a me contentar com o modelo 1.6", diz. No momento, a empresa de Nazarenko assiste uma queda de vendas da Ivana —que teve seu ápice de comercialização no começo deste ano—, mas deve voltar à carga no final de 86. "Estou projetando um equipamento de controle para automação industrial, como trabalho final de curso, que deverá ser o

próximo produto da empresa".

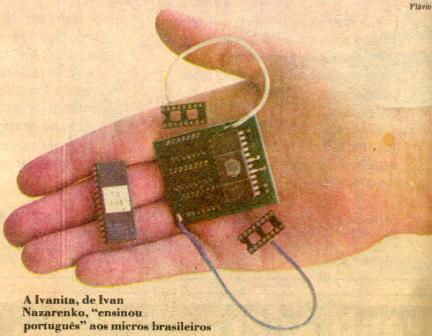
Christovão Manuel Baptista, 32, diretor da Blump, também começa a vislumbrar o retorno de vinte anos investidos em pesquisa. Criando seu primeiro robô aos doze anos —"com lata de talco no corpo e cabeça de tampa de aerosol"— Baptista começa a engatinhar num mercado milionário: o de braços mecânicos e robôs industriais. Dos Cr\$ 50 mil iniciais, abocanhados no programa do Silvio Santos com sua primeira engenhoca, Baptista pode saltar para Cr\$ 100 milhões —o preço estimado de seu modelo mais simples de robô industrial.

Já Fábio Cavalcanti, 24, e Jecel Matos de Assumpção Jr., 24, criaram há dois anos atrás um computador baseado na filosofia educacional Logo, que prevê o uso de micros em salas de aula, batizado de Logo Machine. Na época, quando os micros domésticos ainda não tinham cor, a máquina de Fábio e Jecel, após seis meses de pesquisa, operava com oito cores e sintetizador de sons. Resultado: ninguém quis fabricar a Logo Machine. "Nos antecipamos as necessidades do mercado", reconhece Fábio.

Atualmente, Fábio e Jecel têm um projeto bem mais ambicioso: estão projetando um micro de 32 bits voltado à linguagem Small Talk —uma poderosa filosofia de programação criada pela Xerox dos EUA, que não requer qualquer conhecimento de computação. "Agora estamos no caminho certo", diz Jecel. "Conceitos de Small Talk são usados no Macintosh da Apple, no sistema operacional Unix, mas queremos criar uma máquina puramente Small Talk".



Fábio (esq.), Jecel e a placa de um microcomputador frustrado



A Ivana, de Ivan Nazarenko, "ensinou português" aos micros brasileiros

Figure 46: news item in São Paulo two years later